

BASIC UNDERSTANDING ON CLOUD COMPUTING

Introduction to Cloud Computing

Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction. This cloud model is composed of five essential characteristics, three service models, and four deployment models. The essential characteristics are on-demand self-service, broadcast network access, resource pooling, rapid elasticity, measured service, etc.

Service Models

Software as a Service (SaaS): The capability provided to the consumer is to use the provider's applications running on a cloud infrastructure. The applications are accessible from various client devices through either a thin client interface, such as a web browser (e.g., web-based email), or a program interface. The consumer does not manage or control the underlying cloud infrastructure including network, servers, operating systems, storage, or even individual application capabilities, with the possible exception of limited user specific application configuration settings.

Platform as a Service (PaaS): The capability provided to the consumer is to deploy onto the cloud infrastructure consumer-created or acquired applications created using programming languages, libraries, services, and tools supported by the provider.³ The consumer does not manage or control the underlying cloud infrastructure including network, servers, operating systems, or storage, but has control over the deployed applications and possibly configuration settings for the application-hosting environment.

Infrastructure as a Service (IaaS): The capability provided to the consumer is to provision processing, storage, networks, and other fundamental computing resources where the consumer is able to deploy and run arbitrary software, which can include operating systems and applications. The consumer does not manage or control the underlying cloud infrastructure but has control over operating systems, storage, and deployed applications; and possibly limited control of select networking components (e.g., host firewalls).

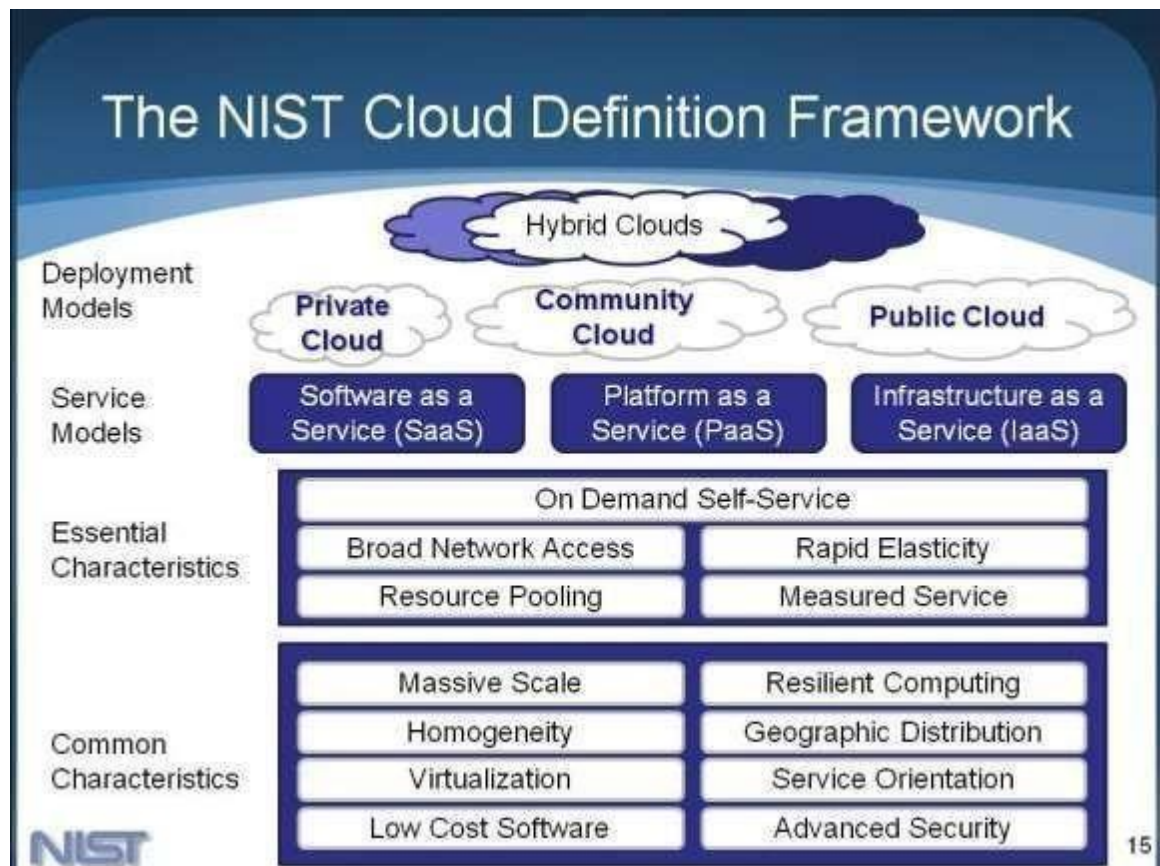
Deployment Models

Private cloud: The cloud infrastructure is provisioned for exclusive use by a single organization comprising multiple consumers (e.g., business units). It may be owned, managed, and operated by the organization, a third party, or some combination of them, and it may exist on or off premises.

Community cloud: The cloud infrastructure is provisioned for exclusive use by a specific community of consumers from organizations that have shared concerns (e.g., mission, security requirements, policy, and compliance considerations). It may be owned, managed, and operated by one or more of the organizations in the community, a third party, or some combination of them, and it may exist on or off premises.

Public cloud: The cloud infrastructure is provisioned for open use by the general public. It may be owned, managed, and operated by a business, academic, or government organization, or some combination of them. It exists on the premises of the cloud provider.

Hybrid cloud: The cloud infrastructure is a composition of two or more distinct cloud infrastructures (private, community, or public) that remain unique entities, but are bound together by standardized or proprietary technology that enables data and application portability (e.g., cloud bursting for load balancing between clouds).



NIST Cloud Definition Framework

EX. NO:

**DATE: CREATION AND MANAGEMENT OF VIRTUAL MACHINE IN
VIRTUALIZED ENVIRONMENT – VMWAREWORKSTATION**

AIM

To install VMWare Workstation and configure a virtual machine.

THEORY

VIRTUALIZATION

Virtualization is creation of Virtual Machines which can emulate hardware in software or in other words it is the creation of virtual version of something such as a hardware platform, operating system, storage device, or network resources (from Wikipedia). Virtualization is achieved or created with the help of software and this particular software allows you to install any number of OS on your system without using the available hardware directly. When you are running an OS over the top of another on your machine the whole environment acts like a HOST and GUEST OS. The real operating system acts as a HOST and the OS run by virtualization software acts as a GUEST OS. The entire load balancing is actually done by the HOST operating system.

Types of Virtualization: There are mainly three types of virtualization.

- Full virtualization
- OS level virtualization
- Para virtualization

Full virtualization

As the name suggests everything in a system is virtualized which includes the processor, storage, networking components etc. Virtual Box, VMware are example of “Full Virtualization” solutions.

OS Level virtualization

In this type of virtualization only applications are run inside the software. In this case the application is given a platform to work. Isolation is created and the application is made to believe that it is the only thing running on the system.

Para virtualization

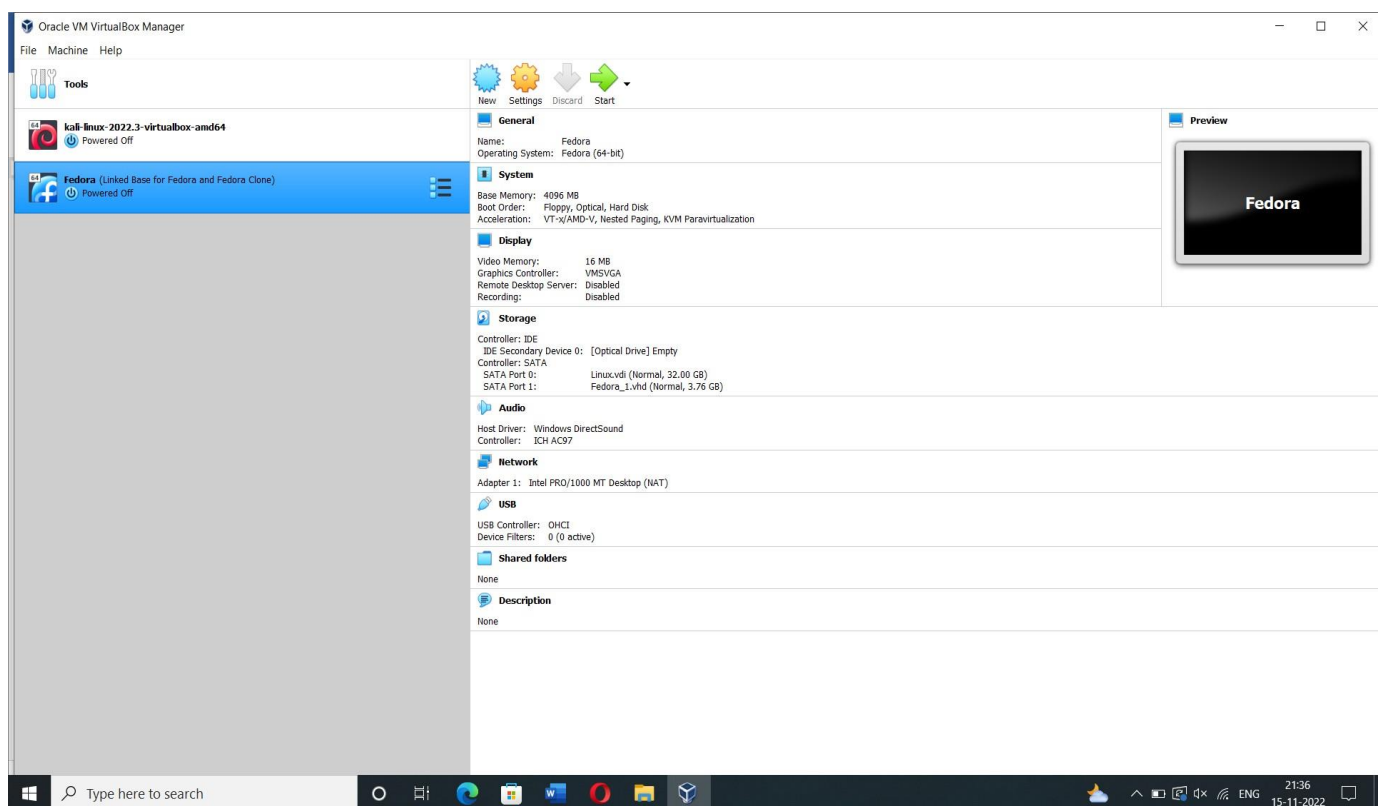
It's a semi-virtualized environment created for the guest OS. A modified guest OS is created using a hypervisor. “The intent of the modified interface is to reduce the portion of the guest's execution time spent performing operations which are substantially more difficult to run in a virtual environment compared to a non-virtualized environment. The Para virtualization provides specially defined „hooks“ to allow the guest(s) and host to request and acknowledge these tasks, which would otherwise be executed in the virtual domain (where execution performance is worse). A successful Para virtualized platform may allow the virtual machine monitor (VMM) to be simpler (by relocating execution of critical tasks from the virtual domain to


the host domain), and/or reduce the overall performance degradation of machine-execution inside the virtual-guest.”

Advantages of Virtualization

- ❑ One of the biggest advantages of virtualization is scalability i.e. the ability to expand. Whenever there is excessive load on some part of application in a server you can easily create a similar virtual environment on a different server and configure the setup.
- Hardware maintenance cost is reduced because you don’t need many servers to install different applications.
- ❑ You can save a huge amount of energy by running one physical server instead of many and less power backup is required.
- ❑ You can get faster and safer backups by taking live snapshot while server is running.
- ❑ You will get centralized monitoring of your resources as virtualization provides easy way of connecting and maintaining your virtual servers

OUTPUT:





RESULT

Thus the procedure to create virtual machine with instance and virtual clone is done successfully.

EX. NO:

DATE:

VIRTUALIZE A MACHINE AND CHECK HOW MANY MACHINES CAN BE UTILIZED AT A PARTICULAR TIME

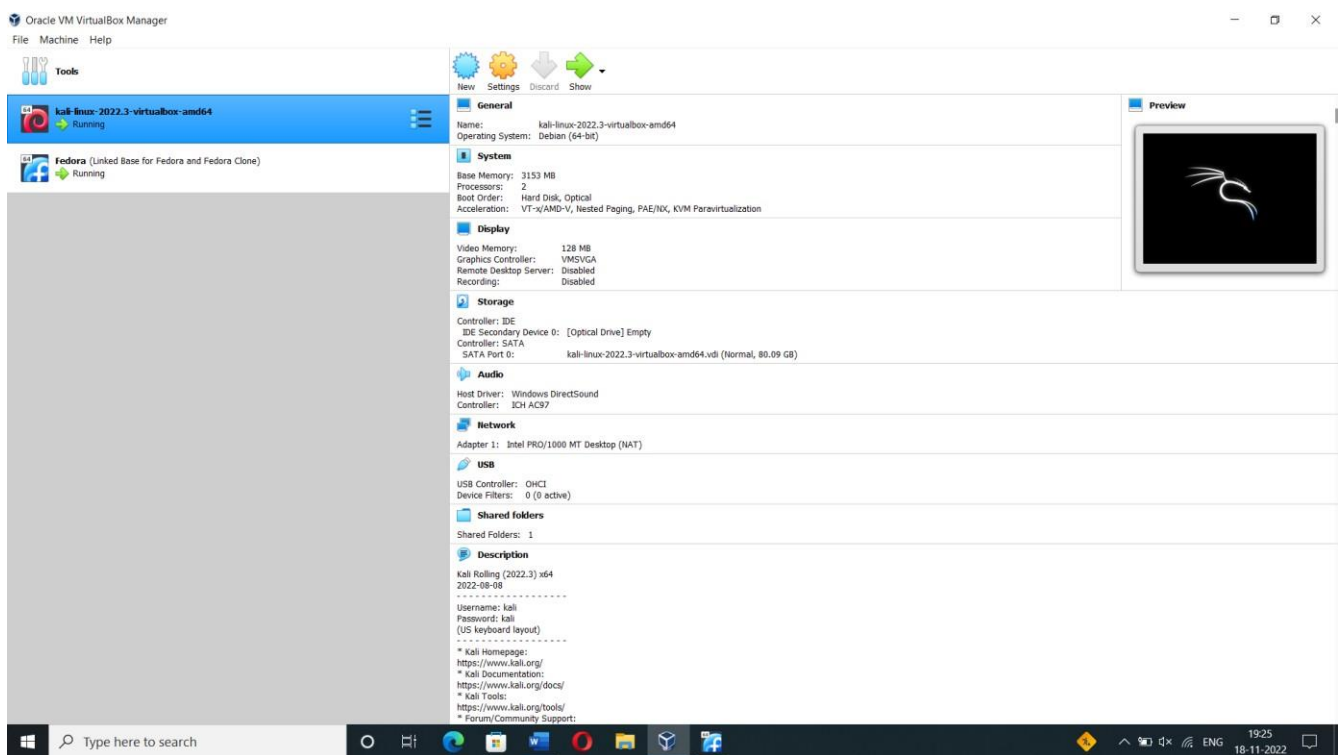
AIM

To virtualize a machine and check how many machines can be utilized at a particular time.

PROCEDURE

1. Open the VMware.
2. Run the virtual machines installed in the VMware.
3. The virtual machines run parallelly in the VMware.
4. We can observe that many virtual machines can be run at a particular time.

OUTPUT



RESULT

Thus virtualize a machine and checking the number of machines running particularly at a time has been done successfully.

EX. NO:

DATE:

CREATE VM CLONE AND ATTACH A VIRTUAL BLOCK

AIM

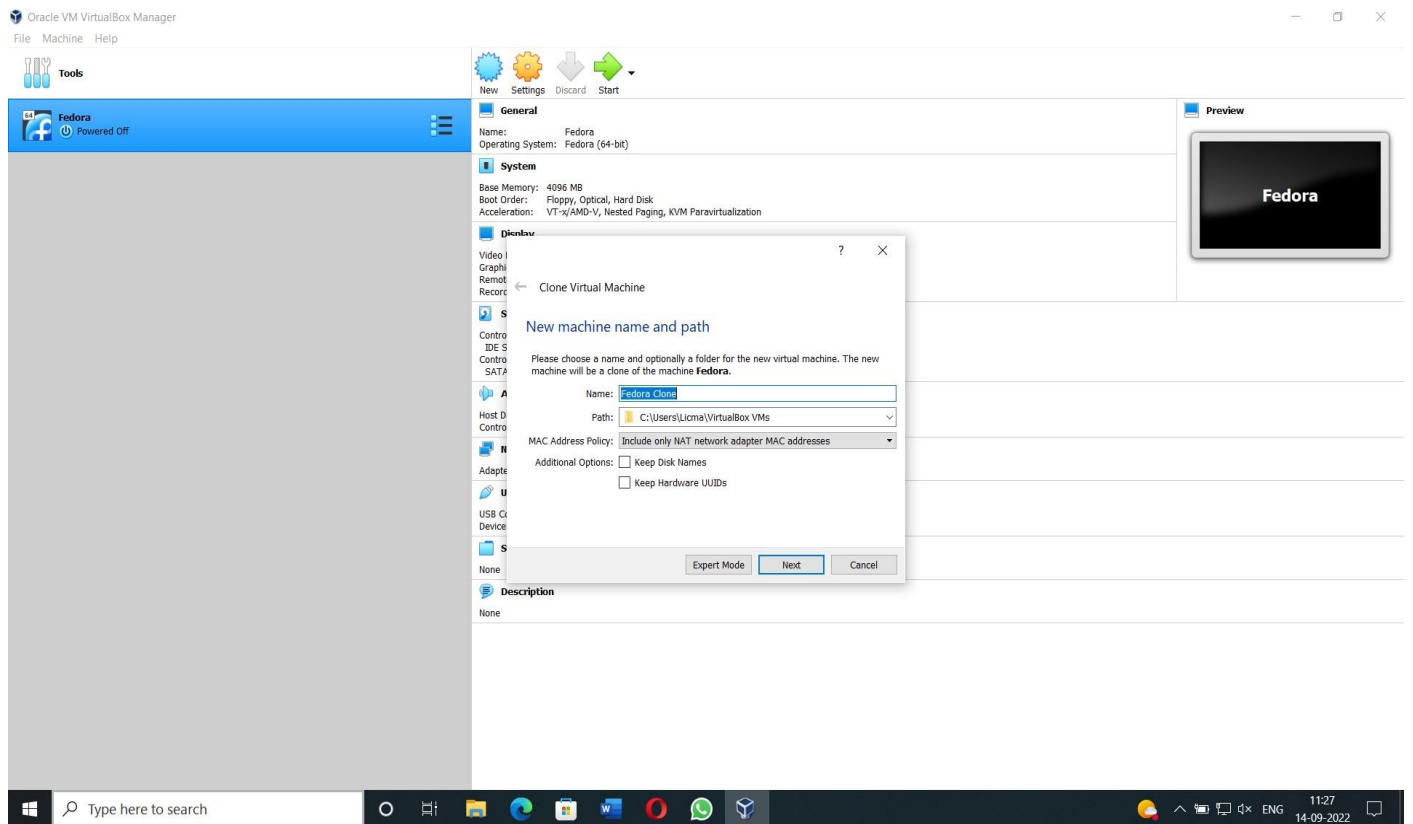
To find procedure to attach a virtual block to virtual machine and check whether it holds data even after the release of the virtual machine.

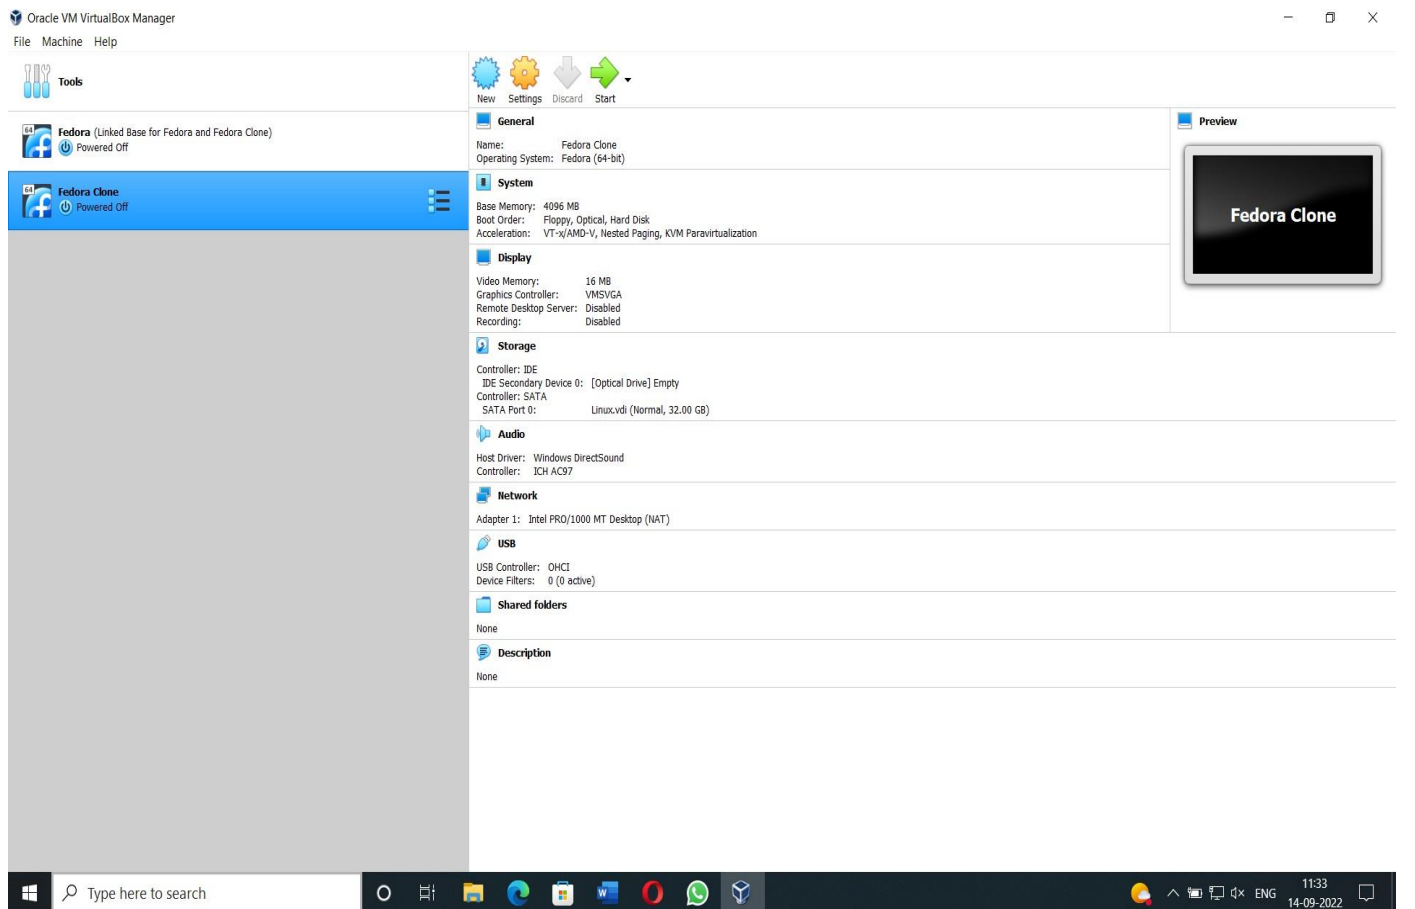
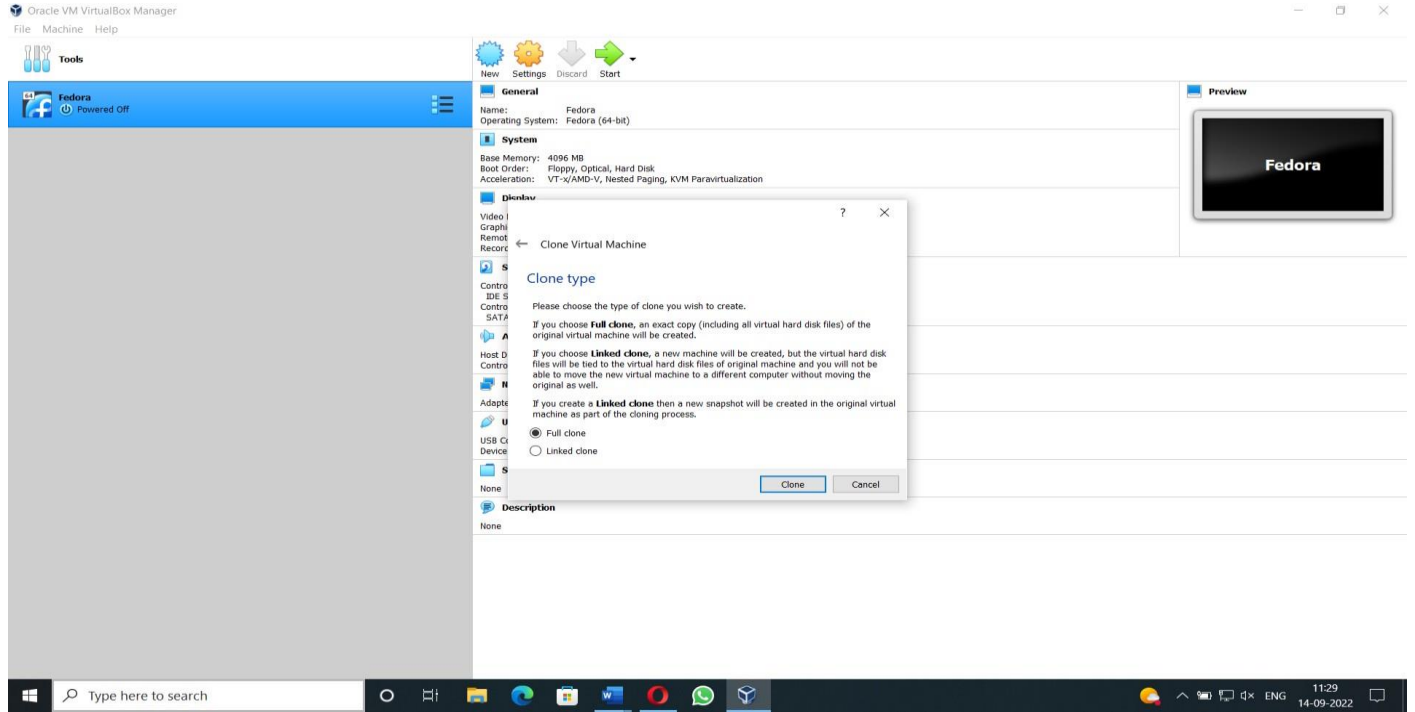
- a) Create a VM clone
- b) Adding an additional block to Hard-Disk in Virtual Machine.

PROCEDURE

a) Create a VM clone

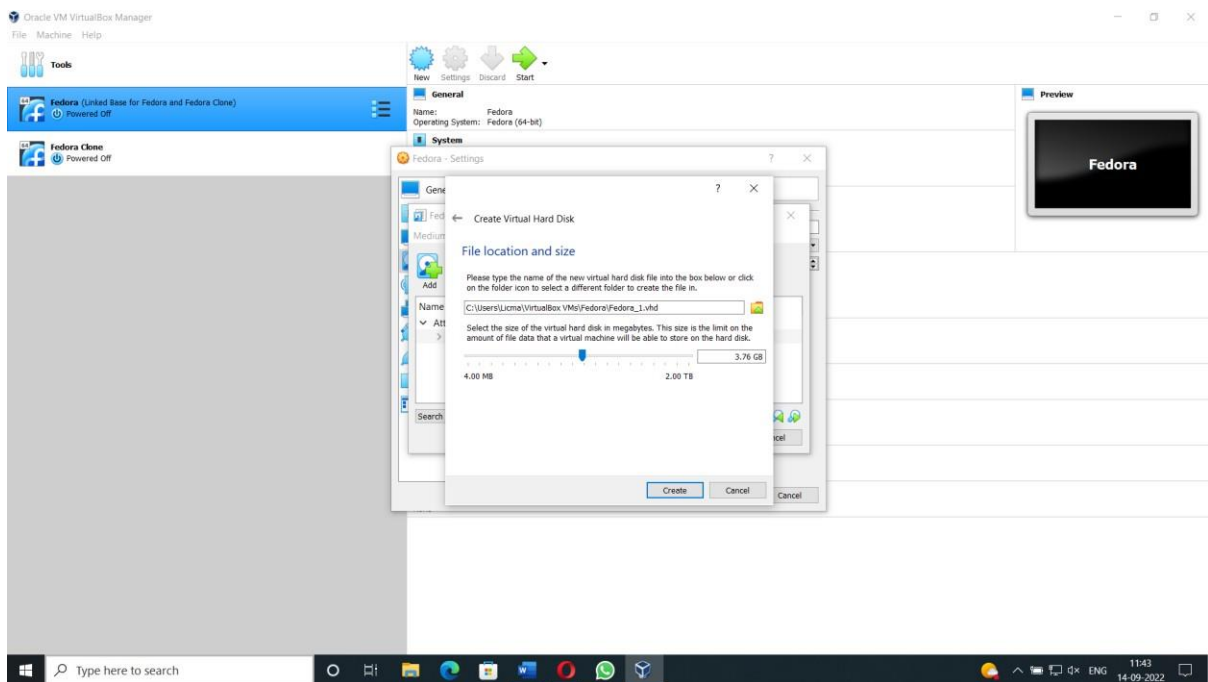
Step: 1 –Click VM -> Manage -> Clone



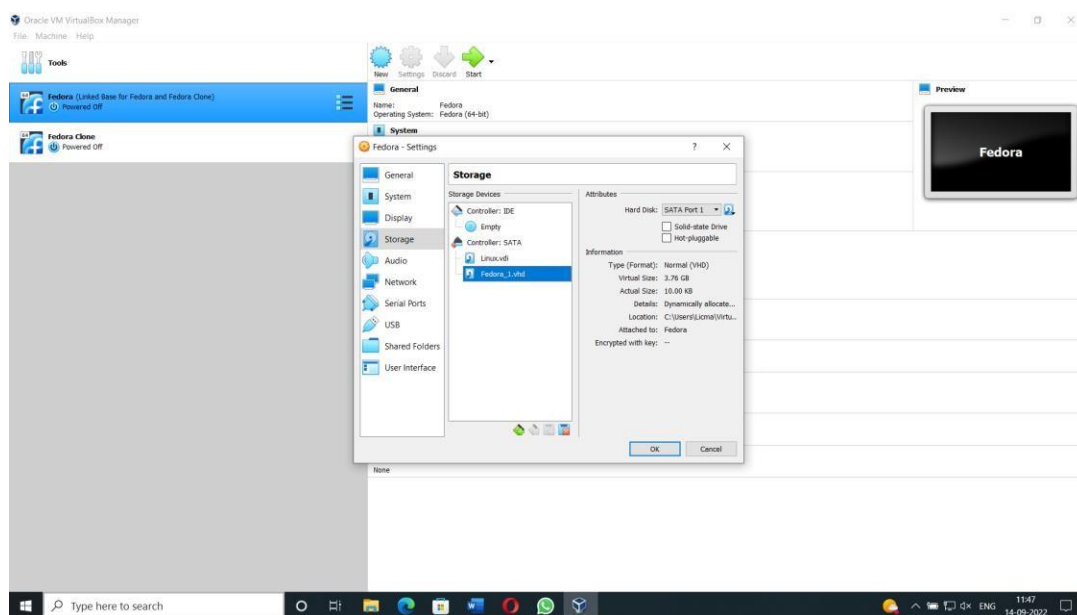


b) Adding an additional block to Hard-Disk in Virtual Machine.

Step: 1 - Select hard Disk from the virtual machine settings dialog box. Select add button. In the Add Hardware wizard select the hardware type as hard disk and click next. Select the virtual disk type as IDE and click next.



Step: 2 – In the next page select the Create a new virtual disk option and click next. Provide the disk capacity as 20GB and select store virtual disk as single file option and click next. Specify the disk file and click finish and virtual block is created.



RESULT

Thus the procedure to create a VM clone and adding an additional block to Hard-Disk in virtual machine is completed successfully.

EX NO:

DATE:

SALESFORCE

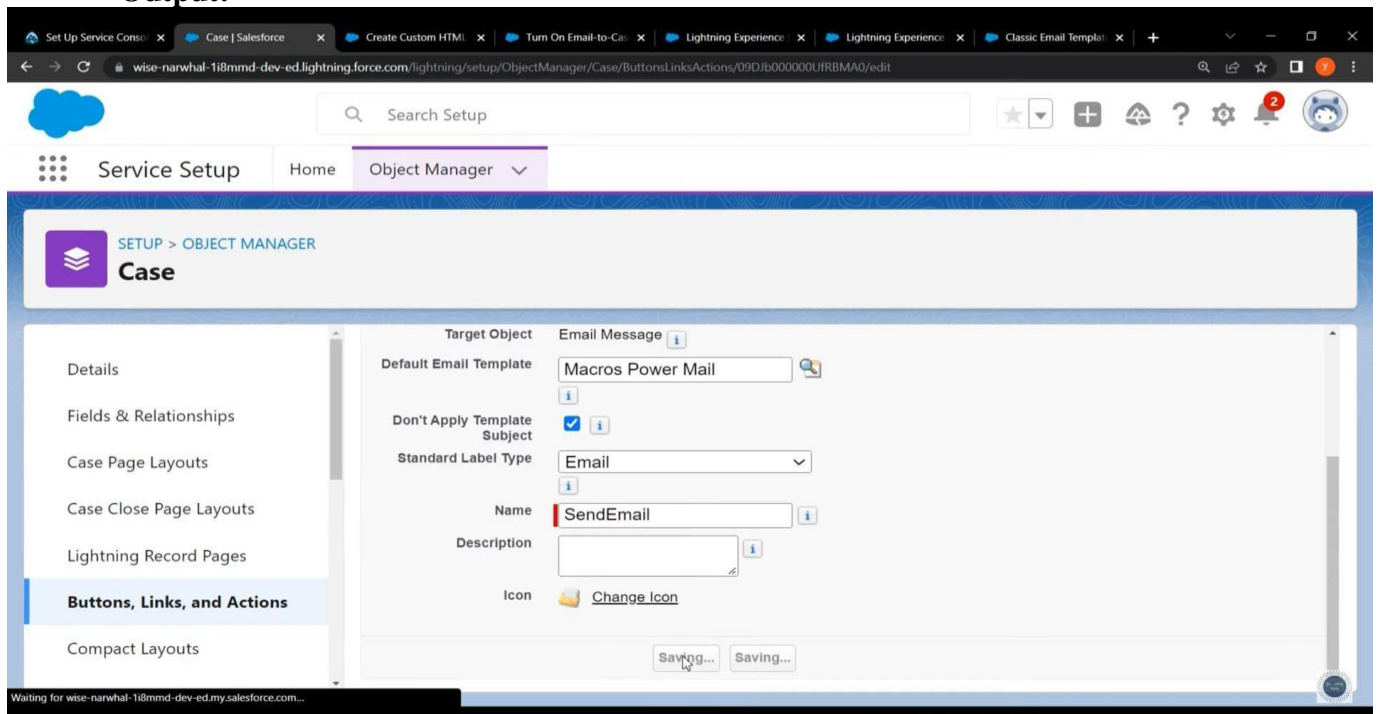
Aim:

To develop a simple E-mail automation service using salesforce.

Procedure :

- 1) Open Salesforce and login to the account.
- 2) Open the Service cloud agent productivity module
- 3) Do the setup service console productivity tools
 - a. Give users access to quick text and macros.
 - b. Customize your app for quick text, macros, and history.
 - c. Set up mass quick actions.
 - d. Customize an email action.
 - e. Enable email notifications for case owners
- 4) Open the Create Macros and Quick Text to Reduce Clicks
 - a. Create macros.
 - b. Create quick text.
- 5) Open the module Use All the Service Console Productivity Tools Together
 - a. Use split view.
 - b. Run a macro.
 - c. Use and find the keyboard shortcuts.
 - d. Insert quick text.
 - e. Perform mass quick actions.
 - f. Use the History utility.

Output:



Service Console Quick Text

Quick Text Recent

1 item

Search recent quick text...

New Quick Text New Folder

QUICK TEXT	Quick Text Name	Message	Folder	Channel	Category
Recent	Solar Panel Apology	{!Case.CaseNumber} Hello		Email	Closings

QUICK TEXT

- Recent
- All Quick Text

FOLDERS

- All Folders
- Created by Me
- Shared with Me

FAVORITES

Macros History

Service Console Cases

Replacement Owner

Case Owner

yourcoddng buddy

Perform action

Update 2 records and change these fields: Internal Comments, Status?

Cancel OK

Fields to update

Internal Comments, Status

Cancel Save

Case Nu...	Contact Name
1 00001002 Stella Pavlova	
2 00001016 Edna Frank	
3 00001024 Lauren Boyle	

Date/Time Opened	Case O...
5/30/2022, 7:47 PM	ybudd
5/30/2022, 7:47 PM	ybudd
5/30/2022, 7:47 PM	ybudd



Module

Service Cloud Agent Productivity

Help your agents work faster and more efficiently in the Service Console.



+1,600 POINTS

Completed 6/14/22

Learn About Service Console Productivity Tools

~10 mins



Set Up Service Console Productivity Tools

~30 mins



Create Macros and Quick Text to Reduce Clicks

~30 mins



Use All the Service Console Productivity Tools Together

~30 mins



RESULT

Thus the procedure to create an email automation service using salesforce is completed successfully.

EX NO:

DATE:

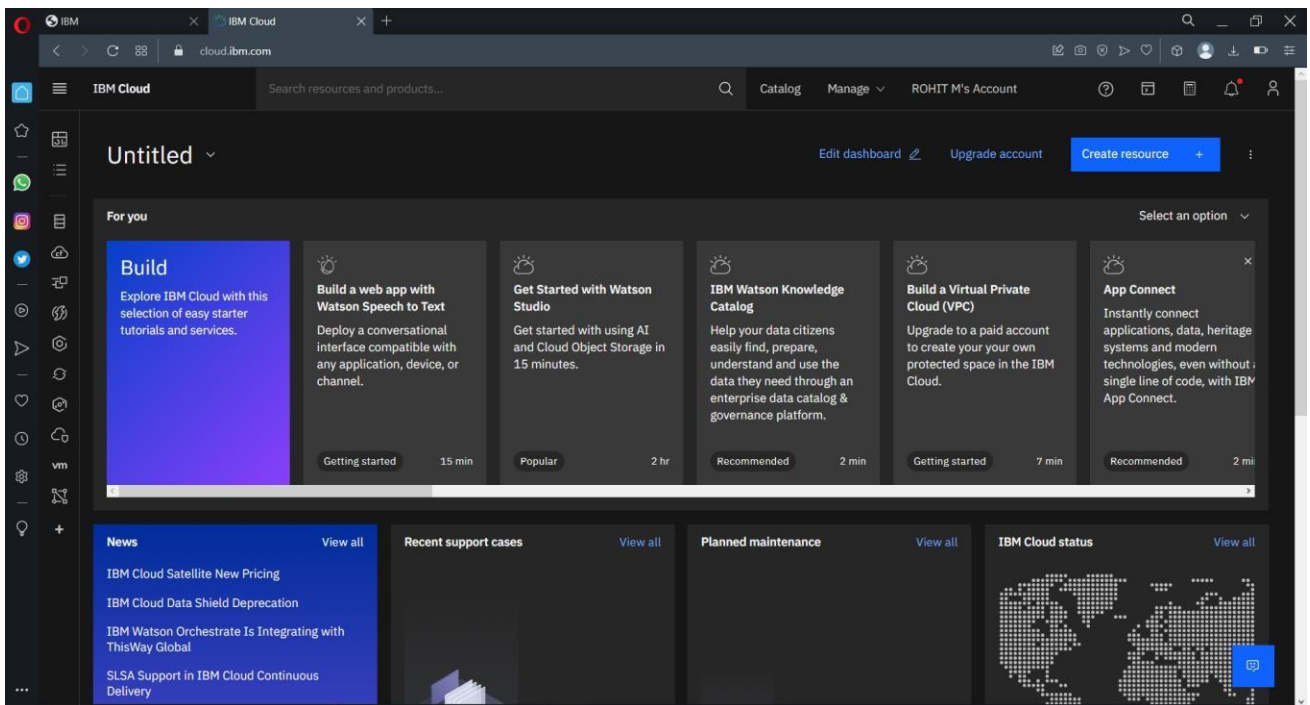
LAUNCH A CLOUD INSTANCE USING IBM CLOUD

Aim:

To launch a cloud instance using a public IaaS cloud service like the IBM cloud.

Procedure :

1. Register to IBM Cloud
2. Sign in with your credentials
3. Log in to your IBM Cloud account, and click on Catalog



4. Type Cloudant in the Search bar and click to open it.
5. Select an offering and an environment
6. Select region as Dallas & Type an instance name then click on create service.

7. After you click create the system displays a message to say that the instance is being provisioned, which returns you to the Resource list. From the Resource list, you see that the status for your instance is, Provision in progress.
8. When the status changes to Active, click the instance.

OUTPUT

Name	Group	Location	Product	Status	Tags
Filter by name or IP address...	Filter by group or org...	Filter...	Filter...	Filter...	Filter...
Compute (0)					
Containers (0)					
Networking (0)					
Storage (0)					
AI / Machine Learning (0)					
Analytics (0)					
Blockchain (0)					
Databases (1)					
Cloudant-gp	Default	Dallas	Cloudant	Active	—
Developer tools (0)					
Logging and monitoring (0)					
Migration (0)					

RESULT

Thus the launching of a cloud instance using a public IaaS cloud service like the IBM cloud has been done successfully.

EX. NO:

DATE:

PUBLIC CLOUD SERVICE

AIM:

To work with a public cloud service such as ServiceNow.

DESCRIPTION

Visual Task Boards (VTB) transform the navigation of lists and forms into an interactive graphical experience. With Visual Task Boards, you can view and update multiple task records, which appear as *cards* that can be moved between *lanes*. An *activity stream* on the board displays recent activity so you can easily track changes to tasks. You can add task cards from any table that extends Task to intuitively and easily track updates and edit records directly from the board. Any user can use task boards, regardless of role, though access control rules (ACLs) may limit which cards each user can see. The Visual Task Board interface provides a graphic-rich environment suited for managing and collaborating on records. For example, a support manager might create a board for the team to track their assigned incidents by state in real time.

OUTPUT

The screenshot displays the ServiceNow Visual Task Board (VTB) interface. The top navigation bar includes the ServiceNow logo, a search bar, and links for Home, Certification, Help, and a user profile. The left sidebar lists various activities and resources, including 'ServiceNow Fundamentals Course Summary' and 'ServiceNow Fundamentals Simulator (San Diego)'. The main area shows a VTB with multiple lanes representing different states or priorities. Cards are visible in the lanes, each containing task details like 'Upgrade NY RAC to Oracle 12c'. Below the VTB, there are instructions for acting as the Change Request coordinator, including assigning tasks to specific users like Bud Richman and Charlie Whitherspoon. The bottom right corner features a chat icon.

2. Acting as the Change Request coordinator, begin to even out the workload using your Visual Task Board as follows:

- Assign the currently unassigned 'Upgrade Tomcat server to 4.0' Change Request to **Bud Richman**.
- Looks like Fred Luddy is swamped. Re-assign his 'Oracle FLX (Java App Dependency)' Change Request to **Charlie Whitherspoon**.

Change Requests by Assigned to						
Guided Board						
<div> <div>Filter by title or number</div> <div>Due By</div> <div> <div></div> <div></div> <div></div> <div></div> </div> </div>						
<div> <div>Beth Anglin</div> <div>Bow Ruggeri</div> <div>Bad Richman</div> <div>Charlie Whitherspoon</div> <div>Christen Mitchell</div> <div>David Loo</div> <div>Don Goodliffe</div> <div>Fred Luddy</div> <div>ITIL User</div> <div>System Administrator</div> </div>						
3	Christen Mitchell 1	David Loo 32	Don Goodliffe 1	Bow Ruggeri 25	Fred Luddy 6	+ Add Lane
her 100 Gb drive on the 2nd FL...	CMS App FLX (depends on Java Apps FLX)	Another Java Application Server change	Java Application Server change	Update /etc/network/interfaces to incl...	Upgrade NY RAC to Oracle 12c	
CHG0000006 7mo ago	CM CHG0000014 7mo ago	CHG0000011 7mo ago	CHG0000010 7mo ago	BR CHG0000093 7mo ago	FL CHG0000088 7mo ago	
ts to know what it'd cost to s...	+ Add Card	A failed RAID controller card needs to b...	+ Add Card	Update /etc/network/interfaces to incl...	Upgrade NY RAC to Oracle 12c	
CHG0000007 7mo ago		CHG0000087 7mo ago		BR CHG0000092 7mo ago	FL CHG0000078 7mo ago	
ew Cisco		Increase db_block_buffers from 5000 t...		Update /etc/network/interfaces to incl...	Upgrade NY RAC to Oracle 12c	
CHG0000008 7mo ago		CHG0000094 7mo ago		BR CHG0000091 7mo ago	FL CHG0000068 7mo ago	
+ Add Card		Deploy new Cisco Catalyst 4500		Update default support group on all Or...	Upgrade NY RAC to Oracle 12c	
		CHG0000089 7mo ago		BR CHG0000090 7mo ago	FL CHG0000058 7mo ago	
		Upgrade OWA-SD-01 to MS Windows Se...		Update /etc/network/interfaces to incl...	Upgrade NY RAC to Oracle 12c	
		CHG0000095 7mo ago		BR CHG0000082 7mo ago	FL CHG0000048 7mo ago	
		Change default router on unix201		Update /etc/network/interfaces to incl...	Upgrade NY RAC to Oracle 12c	
		CHG0000096 7mo ago		BR CHG0000081 7mo ago	FL CHG0000038 7mo ago	
		A failed RAID controller card needs to b...		Update /etc/network/interfaces to incl...	+ Add Card	
		CHG0000077 7mo ago		BR CHG0000083 7mo ago		
		Increase db_block_buffers from 5000 t...		Update default support group on all Or...		
		CHG0000084 7mo ago		BR CHG0000080 7mo ago		
		Deploy new Cisco Catalyst 45...	+ Add Card	Update /etc/network/interfac...	+ Add Card	
		CHG0000079 7mo ago		BR CHG0000073 7mo ago		

RESULT

Thus the procedure to work with public cloud service such as servicenow has been done successfully.

EX. NO:

DATE:

MODEL A CLOUD ENVIRONMENT USING CLOUD SIMULATORS

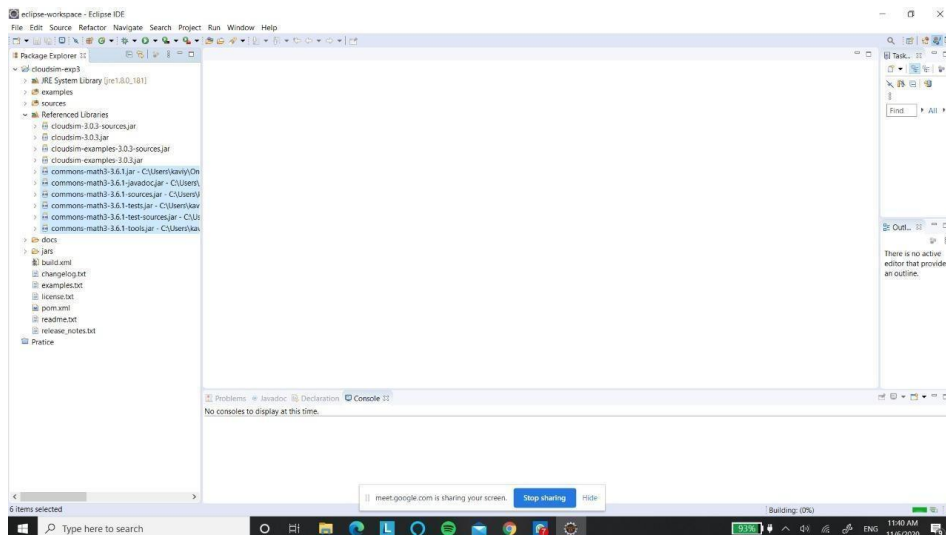
AIM:

Installation of CloudSim into Eclipse to model a cloud computing environment and analyze the VM provisioning using CloudSim.

PROCEDURE: (Installation Steps)

- 1) Open up Eclipse and go to Menu Section, then click File, keep on clicking New and finally select java project.
- 2) A JRE environment is required and hence make sure to have jdk installed and finally give a name to the project.
- 3) CloudSim package is required to be downloaded and extract the CloudSim zip file to set the directory for the project.
- 4) Common Math from apache is a jar file which is also required to run the math functions in the CloudSim file.
- 5) Now all the files are installed in the Eclipse environment.
- 6) Run a program in Cloudsim environment to check if the simulation is complete.

OUTPUT:



RESULT :

Thus program for the given scenario has been compiled and executed successfully.

EX. NO:

DATE:

**IMPLEMENT ROUNDROBIN TASK SCHEDULING IN BOTH TIMESHARED AND
SPACESHARED CPU ASSIGNMENT**

AIM:

To Model a Cloud computing environment having Data centre that had 100 hosts. The hosts are to be modelled to have a CPU core (1000 MIPS), 2 GB of RAM and 1 TB of storage. Consider the workload model for this evaluation including provisioning requests for 400 VMs, with each request demanding 1 CPU core (250 MIPS), 256 MB of RAM and 1 GB of storage. Each VM hosts a web-hosting application service, whose CPU utilization distribution was generated according to the uniform distribution. Each instance of a webhosting service required 150,000 MIPS or about 10 minutes to complete execution assuming 100% utilization.

PROCEDURE :

- 1) Initialize the CloudSim library and datacenters.
- 2) Create 2 VM's and the second VM will have twice the priority of VM1 and so will receive twice CPU time. Add the VMs to the vmList.
- 3) Create two Cloudlets. Where,
Cloudlet properties
 intid = 0;
 longlength = 40000;
 longfileSize = 300;
 longoutputSize = 300;
- 4) Add the cloudlets to a list and submit cloudlet list to the broker.
- 5) Bind the cloudlets to the vms. This way, the broker will submit the bound cloudlets only to the specific VM.
- 6) Create a datacenter. The steps needed to create a PowerDatacenter,
 - (i) We need to create a list to store our machine
 - (ii) Machine contains one or more PEs or CPUs/Cores.
 - (iii) Create PEs and add these into a list.
 - (iv) Create Hosts with its id and list of PEs and add them to the list of machines
 - (v) create another machine in the Data center.
 - (vi) Create a DatacenterCharacteristics object and create a PowerDatacenter object.
- 7) Print the Cloudlet objects and observe the result

PROGRAM FOR SPACE SHARED :-

```
package org.cloudbus.cloudsim.examples;

import java.text.DecimalFormat;
import java.util.ArrayList;
import java.util.Calendar;
import java.util.LinkedList;
import java.util.List;
import org.cloudbus.cloudsim.Cloudlet;
import org.cloudbus.cloudsim.CloudletSchedulerSpaceShared;
import org.cloudbus.cloudsim.Datacenter;
import org.cloudbus.cloudsim.DatacenterBroker;
import org.cloudbus.cloudsim.DatacenterCharacteristics;
import org.cloudbus.cloudsim.Host;
import org.cloudbus.cloudsim.Log;
import org.cloudbus.cloudsim.Pe;
import org.cloudbus.cloudsim.Storage;
import org.cloudbus.cloudsim.UtilizationModel;
import org.cloudbus.cloudsim.UtilizationModelFull;
import org.cloudbus.cloudsim.Vm;
import org.cloudbus.cloudsim.VmAllocationPolicySimple;
import org.cloudbus.cloudsim.VmSchedulerTimeShared;
import org.cloudbus.cloudsim.core.CloudSim;
import org.cloudbus.cloudsim.provisioners.BwProvisionerSimple;
import org.cloudbus.cloudsim.provisioners.PeProvisionerSimple;
import org.cloudbus.cloudsim.provisioners.RamProvisionerSimple;

publicclass CloudSimExample3 {

    /** The cloudlet list. */
    privatestatic List<Cloudlet>cloudletList;

    /** The vmlist. */
    privatestatic List<Vm>vmList;

    /**
     * Creates main() to run this example
     */
    publicstaticvoid main(String[] args) {

        Log.println("Starting CloudSimExample3...");

        try {
            // First step: Initialize the CloudSim package. It should be called
            // before creating any entities.
            intnum_user = 1; // number of cloud users
            Calendar calendar = Calendar.getInstance();
            booleantrace_flag = false; // mean trace events
```

```

// Initialize the CloudSim library
CloudSim.init(num_user, calendar, trace_flag);

// Second step: Create Datacenters
//Datacenters are the resource providers in CloudSim. We need at list one of
them to run a CloudSim simulation
@SuppressWarnings("unused")
Datacenter datacenter0 = createDatacenter("Datacenter_0");
DatacenterBroker broker = createBroker();
intbrokerId = broker.getId();

//Fourth step: Create one virtual machine
vmlist = new ArrayList<Vm>();

//VM description
intvmid = 0;
intmips = 250;
longsize = 10000; //image size (MB)
intram = 2048; //vm memory (MB)
longbw = 1000;
intpesNumber = 1; //number of cpus
String vmm = "Xen"; //VMM name

//create two VMs
Vm vm1 = new Vm(vmid, brokerId, mips, pesNumber, ram, bw, size, vmm,
new CloudletSchedulerSpaceShared());

//the second VM will have twice the priority of VM1 and so will receive twice
CPU time
vmid++;
Vm vm2 = new Vm(vmid, brokerId, mips * 2, pesNumber, ram, bw, size, vmm,
new CloudletSchedulerSpaceShared());

//add the VMs to the vmList
vmlist.add(vm1);
vmlist.add(vm2);

// add the VMs to the vmList
broker.submitVmList(vmlist);

//Fifth step: Create two Cloudlets
cloudletList = new ArrayList<Cloudlet>();

//Cloudlet properties
intid = 0;
longlength = 40000;
longfileSize = 300;
longoutputSize = 300;
UtilizationModel utilizationModel = new UtilizationModelFull();

```

```

        Cloudlet cloudlet1 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet1.setUserId(brokerId);

        id++;
        Cloudlet cloudlet2 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet2.setUserId(brokerId);
        id++;
        Cloudlet cloudlet3 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet3.setUserId(brokerId);

        id++;
        Cloudlet cloudlet4 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet4.setUserId(brokerId);

        id++;
        Cloudlet cloudlet5 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet5.setUserId(brokerId);

        id++;
        Cloudlet cloudlet6 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet6.setUserId(brokerId);

        id++;
        Cloudlet cloudlet7 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet7.setUserId(brokerId);

        id++;
        Cloudlet cloudlet8 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet8.setUserId(brokerId);

        id++;

        //add the cloudlets to the list
        cloudletList.add(cloudlet1);
        cloudletList.add(cloudlet2);
        cloudletList.add(cloudlet3);
        cloudletList.add(cloudlet4);
        cloudletList.add(cloudlet5);
        cloudletList.add(cloudlet6);
        cloudletList.add(cloudlet7);
        cloudletList.add(cloudlet8);

        //submit cloudlet list to the broker
        broker.submitCloudletList(cloudletList);

```

```

        //bind the cloudlets to the yms. This way, the broker
        // will submit the bound cloudlets only to the specific VM
        broker.bindCloudletToVm(cloudlet1.getId(),vm1.getId());
        broker.bindCloudletToVm(cloudlet2.getId(),vm1.getId());
        broker.bindCloudletToVm(cloudlet3.getId(),vm1.getId());
        broker.bindCloudletToVm(cloudlet4.getId(),vm1.getId());
        broker.bindCloudletToVm(cloudlet5.getId(),vm2.getId());
        broker.bindCloudletToVm(cloudlet6.getId(),vm2.getId());
        broker.bindCloudletToVm(cloudlet7.getId(),vm2.getId());
        broker.bindCloudletToVm(cloudlet8.getId(),vm2.getId());

        // Sixth step: Starts the simulation
        CloudSim.startSimulation();

        // Final step: Print results when simulation is over
        List<Cloudlet>newList = broker.getCloudletReceivedList();

        CloudSim.stopSimulation();

    printCloudletList(newList);

        Log.println("CloudSimExample3 finished!");
    }
    catch (Exception e) {
        e.printStackTrace();
        Log.println("The simulation has been terminated due to an unexpected
error");
    }
}

private static Datacenter createDatacenter(String name){

    // Here are the steps needed to create a PowerDatacenter:
    // 1. We need to create a list to store
    //    our machine
    List<Host>hostList = new ArrayList<Host>();

    // 2. A Machine contains one or more PEs or CPUs/Cores.
    // In this example, it will have only one core.
    List<Pe>peList = new ArrayList<Pe>();

    int mips = 1000;

    // 3. Create PEs and add these into a list.
    peList.add(new Pe(0, new PeProvisionerSimple(mips))); // need to store Pe id and
MIPS Rating

    //4. Create Hosts with its id and list of PEs and add them to the list of machines
    int hostId=0;

```

```

intram = 2048; //host memory (MB)
longstorage = 1000000; //host storage
intbw = 10000;

hostList.add(
    new Host(
        hostId,
        new RamProvisionerSimple(ram),
        new BwProvisionerSimple(bw),
        storage,
        peList,
        new VmSchedulerTimeShared(peList)
    )
); // This is our first machine

//create another machine in the Data center
List<Pe>peList2 = new ArrayList<Pe>();

peList2.add(new Pe(0, new PeProvisionerSimple(mips)));

hostId++;

hostList.add(
    new Host(
        hostId,
        new RamProvisionerSimple(ram),
        new BwProvisionerSimple(bw),
        storage,
        peList2,
        new VmSchedulerTimeShared(peList2)
    )
); // This is our second machine


// 5. Create a DatacenterCharacteristics object that stores the
//   properties of a data center: architecture, OS, list of
//   Machines, allocation policy: time- or space-shared, time zone
//   and its price (G$/Pe time unit).
String arch = "x86"; // system architecture
String os = "Linux"; // operating system
String vmm = "Xen";
doubletime_zone = 10.0; // time zone this resource located
doublecost = 3.0; // the cost of using processing in this resource
doublecostPerMem = 0.05; // the cost of using memory in this resource
doublecostPerStorage = 0.001; // the cost of using storage in this resource
doublecostPerBw = 0.0; // the cost of using bw in this resource
LinkedList<Storage>storageList = new LinkedList<Storage>(); //we are not adding
SAN devices by now

DatacenterCharacteristics characteristics = new DatacenterCharacteristics(
arch, os, vmm, hostList, time_zone, cost, costPerMem, costPerStorage, costPerBw);

```

```

// 6. Finally, we need to create a PowerDatacenter object.
Datacenter datacenter = null;
try {
    datacenter = new Datacenter(name, characteristics, new
VmAllocationPolicySimple(hostList), storageList, 0);
} catch (Exception e) {
    e.printStackTrace();
}

returndatacenter;
}

//We strongly encourage users to develop their own broker policies, to submit vms and
cloudlets according
//to the specific rules of the simulated scenario
private static DatacenterBroker createBroker(){

    DatacenterBroker broker = null;
    try {
        broker = new DatacenterBroker("Broker");
    } catch (Exception e) {
        e.printStackTrace();
        return null;
    }
    returnbroker;
}

/**
 * Prints the Cloudlet objects
 * @param list list of Cloudlets
 */
private static void printCloudletList(List<Cloudlet>list) {
    intsize = list.size();
    Cloudlet cloudlet;

    String indent = "  ";
    Log.println();
    Log.println("===== OUTPUT =====");
    Log.println("Cloudlet ID" + indent + "STATUS" + indent +
        "Data center ID" + indent + "VM ID" + indent + "Time" + indent +
        "Start Time" + indent + "Finish Time");

    DecimalFormat dft = new DecimalFormat("###.##");
    for (inti = 0; i<size; i++) {
        cloudlet = list.get(i);
        Log.print(indent + cloudlet.getCloudletId() + indent + indent);

        if (cloudlet.getCloudletStatus() == Cloudlet.SUCCESS){
            Log.print("SUCCESS");

```

```

        Log.println( indent + indent + cloudlet.getResourceId() + indent +
indent + indent + cloudlet.getVmId() +
        indent + indent +
dft.format(cloudlet.getActualCPUTime()) + indent + indent +
dft.format(cloudlet.getExecStartTime()+
        indent + indent + dft.format(cloudlet.getFinishTime()));
    }
}
}
}

```

OUTPUT :-

The screenshot shows the Eclipse IDE interface with the 'Console' window open. The console output displays the simulation progress, including messages from the Broker and Datacenter. A table titled 'OUTPUT' is visible, showing the execution details for 8 cloudlets.

Cloudlet ID	STATUS	Data center ID	VM ID	Time	Start Time	Finish Time
4	SUCCESS	2	1	80	0.1	80.1
0	SUCCESS	2	0	160	0.1	160.1
5	SUCCESS	2	1	80	80.1	160.1
6	SUCCESS	2	1	80	160.1	240.1
1	SUCCESS	2	0	160	160.1	320.1
7	SUCCESS	2	1	80	240.1	320.1
2	SUCCESS	2	0	160	320.1	480.1
3	SUCCESS	2	0	160	480.1	640.1

PROGRAM FOR TIME SHARED :-

```

package org.cloudbus.cloudsim.examples;

import java.text.DecimalFormat;
import java.util.ArrayList;
import java.util.Calendar;
import java.util.LinkedList;
import java.util.List;

import org.cloudbus.cloudsim.Cloudlet;
import org.cloudbus.cloudsim.CloudletSchedulerSpaceShared;
import org.cloudbus.cloudsim.Datacenter;
import org.cloudbus.cloudsim.DatacenterBroker;
import org.cloudbus.cloudsim.DatacenterCharacteristics;
import org.cloudbus.cloudsim.Host;
import org.cloudbus.cloudsim.Log;
import org.cloudbus.cloudsim.Pe;

```

```

import org.cloudbus.cloudsim.Storage;
import org.cloudbus.cloudsim.UtilizationModel;
import org.cloudbus.cloudsim.UtilizationModelFull;
import org.cloudbus.cloudsim.Vm;
import org.cloudbus.cloudsim.VmAllocationPolicySimple;
import org.cloudbus.cloudsim.VmSchedulerTimeShared;
import org.cloudbus.cloudsim.core.CloudSim;
import org.cloudbus.cloudsim.provisioners.BwProvisionerSimple;
import org.cloudbus.cloudsim.provisioners.PeProvisionerSimple;
import org.cloudbus.cloudsim.provisioners.RamProvisionerSimple;

publicclass CloudSimExample3 {

    /** The cloudlet list. */
    privatestatic List<Cloudlet>cloudletList;

    /** The vmlist. */
    privatestatic List<Vm>vmList;

    /**
     * Creates main() to run this example
     */
    publicstaticvoid main(String[] args) {

        Log.println("Starting CloudSimExample3...");

        try {
            // First step: Initialize the CloudSim package. It should be called
            // before creating any entities.
            intnum_user = 1; // number of cloud users
            Calendar calendar = Calendar.getInstance();
            booleantrace_flag = false; // mean trace events

            // Initialize the CloudSim library
            CloudSim.init(num_user, calendar, trace_flag);

            // Second step: Create Datacenters
            // Datacenters are the resource providers in CloudSim. We need at list one of
            them to run a CloudSim simulation
            @SuppressWarnings("unused")
            Datacenter datacenter0 = createDatacenter("Datacenter_0");
            DatacenterBroker broker = createBroker();
            intbrokerId = broker.getId();

            //Fourth step: Create one virtual machine
            vmList = new ArrayList<Vm>();

            //VM description
            intvmid = 0;
            inttmips = 250;
            longsize = 10000; //image size (MB)
            inttram = 2048; //vm memory (MB)

```



```

    longbw = 1000;
    intpesNumber = 1; //number of cpus
    String vmm = "Xen"; //VMM name

    //create two VMs
    Vm vm1 = new Vm(vmid, brokerId, mips, pesNumber, ram, bw, size, vmm,
new CloudletSchedulerSpaceShared());

    //the second VM will have twice the priority of VM1 and so will receive twice
CPU time
    vmid++;
    Vm vm2 = new Vm(vmid, brokerId, mips * 2, pesNumber, ram, bw, size, vmm,
new CloudletSchedulerSpaceShared());

    //add the VMs to the vmList
    vmList.add(vm1);
    vmList.add(vm2);

    // add the VMs to the vmList
    broker.submitVmList(vmList);

    //Fifth step: Create two Cloudlets
    cloudletList = new ArrayList<Cloudlet>();

    //Cloudlet properties
    intid = 0;
    longlength = 40000;
    longfileSize = 300;
    longoutputSize = 300;
    UtilizationModel utilizationModel = new UtilizationModelFull();

    Cloudlet cloudlet1 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
    cloudlet1.setUserId(brokerId);

    id++;
    Cloudlet cloudlet2 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
    cloudlet2.setUserId(brokerId);
    id++;
    Cloudlet cloudlet3 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
    cloudlet3.setUserId(brokerId);

    id++;
    Cloudlet cloudlet4 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
    cloudlet4.setUserId(brokerId);

    id++;

```

```

        Cloudlet cloudlet5 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet5.setUserId(brokerId);

        id++;
        Cloudlet cloudlet6 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet6.setUserId(brokerId);

        id++;
        Cloudlet cloudlet7 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet7.setUserId(brokerId);

        id++;
        Cloudlet cloudlet8 = new Cloudlet(id, length, pesNumber, fileSize, outputSize,
utilizationModel, utilizationModel, utilizationModel);
        cloudlet8.setUserId(brokerId);

        id++;

        //add the cloudlets to the list
        cloudletList.add(cloudlet1);
        cloudletList.add(cloudlet2);
        cloudletList.add(cloudlet3);
        cloudletList.add(cloudlet4);
        cloudletList.add(cloudlet5);
        cloudletList.add(cloudlet6);
        cloudletList.add(cloudlet7);
        cloudletList.add(cloudlet8);

        //submit cloudlet list to the broker
        broker.submitCloudletList(cloudletList);

        //bind the cloudlets to the vms. This way, the broker
        // will submit the bound cloudlets only to the specific VM
        broker.bindCloudletToVm(cloudlet1.getCloudletId(),vm1.getId());
        broker.bindCloudletToVm(cloudlet2.getCloudletId(),vm1.getId());
        broker.bindCloudletToVm(cloudlet3.getCloudletId(),vm1.getId());
        broker.bindCloudletToVm(cloudlet4.getCloudletId(),vm1.getId());
        broker.bindCloudletToVm(cloudlet5.getCloudletId(),vm2.getId());
        broker.bindCloudletToVm(cloudlet6.getCloudletId(),vm2.getId());
        broker.bindCloudletToVm(cloudlet7.getCloudletId(),vm2.getId());
        broker.bindCloudletToVm(cloudlet8.getCloudletId(),vm2.getId());

        CloudSim.startSimulation();

        List<Cloudlet>newList = broker.getCloudletReceivedList();

```

```

        CloudSim.stopSimulation();

    printCloudletList(newList);

        Log.println("CloudSimExample3 finished!");
    }
    catch (Exception e) {
        e.printStackTrace();
        Log.println("The simulation has been terminated due to an unexpected
error");
    }
}

private static Datacenter createDatacenter(String name){

    List<Host>hostList = new ArrayList<Host>();

    List<Pe>peList = new ArrayList<Pe>();

    int mips = 1000;

    peList.add(new Pe(0, new PeProvisionerSimple(mips
    int hostId=0;
    int ram = 2048; //host memory (MB)
    long storage = 1000000; //host storage
    int bw = 10000;

    hostList.add(
        new Host(
            hostId,
            new RamProvisionerSimple(ram),
            new BwProvisionerSimple(bw),
            storage,
            peList,
            new VmSchedulerTimeShared(peList)
        )
    );
    List<Pe>peList2 = new ArrayList<Pe>();

    peList2.add(new Pe(0, new PeProvisionerSimple(mips)));

    hostId++;

    hostList.add(
        new Host(
            hostId,
            new RamProvisionerSimple(ram),
            new BwProvisionerSimple(bw),

```

```

        storage,
        peList2,
        new VmSchedulerTimeShared(peList2)
    )
);
String arch = "x86"; // system architecture
String os = "Linux"; // operating system
String vmm = "Xen";
double time_zone = 10.0; // time zone this resource located
double cost = 3.0; // the cost of using processing in this resource
double costPerMem = 0.05; // the cost of using memory in this resource
double costPerStorage = 0.001; // the cost of using storage in this resource
double costPerBw = 0.0; // the cost of using bw in this resource
LinkedList<Storage>storageList = new LinkedList<Storage>(); //we are not adding
SAN devices by now

```

```

DatacenterCharacteristics characteristics = new DatacenterCharacteristics(
arch, os, vmm, hostList, time_zone, cost, costPerMem, costPerStorage, costPerBw);

```

```

// 6. Finally, we need to create a PowerDatacenter object.
Datacenter datacenter = null;
try {
    datacenter = new Datacenter(name, characteristics, new
VmAllocationPolicySimple(hostList, storageList, 0);
} catch (Exception e) {
    e.printStackTrace();
}

return datacenter;
}

```

```

//We strongly encourage users to develop their own broker policies, to submit vms and
cloudlets according

```

```

//to the specific rules of the simulated scenario
private static DatacenterBroker createBroker(){

```

```

    DatacenterBroker broker = null;
    try {
        broker = new DatacenterBroker("Broker");
    } catch (Exception e) {
        e.printStackTrace();
        return null;
    }
    return broker;
}

```

```

/**

```

```

 * Prints the Cloudlet objects
 * @param list list of Cloudlets
 */

```

```

private static void printCloudletList(List<Cloudlet>list) {
    int size = list.size();
}

```

```

Cloudlet cloudlet;

String indent = "  ";
Log.println();
Log.println("===== OUTPUT =====");
Log.println("Cloudlet ID" + indent + "STATUS" + indent +
            "Data center ID" + indent + "VM ID" + indent + "Time" + indent +
"Start Time" + indent + "Finish Time");


DecimalFormat dft = new DecimalFormat("###.##");
for (inti = 0; i<size; i++) {
    cloudlet = list.get(i);
    Log.print(indent + cloudlet.getCloudletId() + indent + indent);

    if (cloudlet.getCloudletStatus() == Cloudlet.SUCCESS){
        Log.print("SUCCESS");

        Log.println( indent + indent + cloudlet.getResourceId() + indent +
indent + indent + cloudlet.getVmId() +
                    indent + indent +
dft.format(cloudlet.getActualCPUTime()) + indent + indent +
dft.format(cloudlet.getExecStartTime())+
                    indent + indent + dft.format(cloudlet.getFinishTime()));
    }
}
}
}

```

OUTPUT :



```

//Submit vm task to the broker
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000

```

RESULT:

Thus program for the given scenario has been compiled and executed successfully for space and time shared.

EX. NO:

DATE: **HADOOP INSTALLATION**

AIM

To set-up one node Hadoop cluster

- System Update
- Install Java
- Add a dedicated Hadoop user
- Install SSH and setup SSH certificates
- Check if SSH works
- Install Hadoop
- Modify Hadoop config files
- Format Hadoop filesystem
- Start Hadoop
- Check Hadoop through web UI
- Stop Hadoop

THEORY

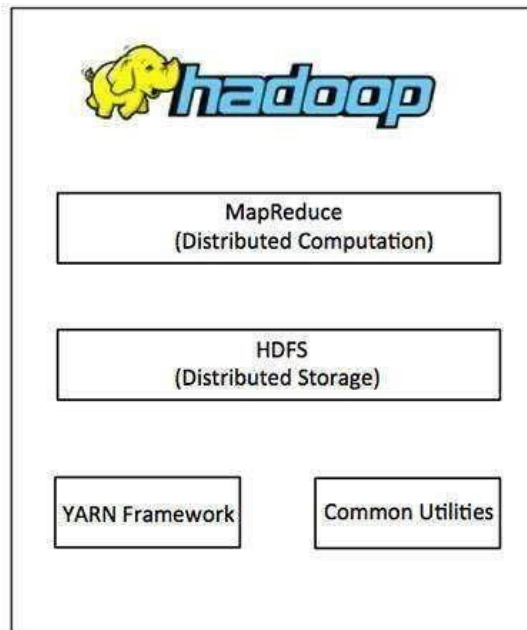
Hadoop is an Apache open source framework written in java that allows distributed processing of large datasets across clusters of computers using simple programming models. A Hadoop frame-worked application works in an environment that provides distributed storage and computation across clusters of computers. Hadoop is designed to scale up from single server to thousands of machines, each offering local computation and storage.

HADOOP ARCHITECTURE

Hadoop framework includes following four modules:

- ▢ **Hadoop Common:** These are Java libraries and utilities required by other Hadoop modules. These libraries provides filesystem and OS level abstractions and contains the necessary Java files and scripts required to start Hadoop.
- ▢ **Hadoop YARN:** This is a framework for job scheduling and cluster resource management.
- **Hadoop Distributed File System (HDFS™):** A distributed file system that provides high-throughput access to application data.
- ▢ **Hadoop MapReduce:** This is YARN-based system for parallel processing of large data sets.

We can use following diagram to depict these four components available in Hadoop framework.



Hadoop Architecture

Since 2012, the term "Hadoop" often refers not just to the base modules mentioned above but also to the collection of additional software packages that can be installed on top of or alongside Hadoop, such as Apache Pig, Apache Hive, Apache HBase, Apache Spark etc.

PROCEDURE

Step 1 – System Update

\$ sudo apt-get update

```
ubuntu@ubuntu-VirtualBox: ~  
ubuntu@ubuntu-VirtualBox:~$ sudo apt-get update  
[sudo] password for ubuntu:  
Hit http://in.archive.ubuntu.com wily InRelease  
Get:1 http://security.ubuntu.com wily-security InRelease [65.9 kB]  
Get:2 http://in.archive.ubuntu.com wily-updates InRelease [65.9 kB]  
Get:3 http://security.ubuntu.com wily-security/main Sources [53.8 kB]  
Hit http://in.archive.ubuntu.com wily-backports InRelease  
Get:4 http://security.ubuntu.com wily-security/restricted Sources [2,854 B]  
Get:5 http://security.ubuntu.com wily-security/universe Sources [13.9 kB]  
Get:6 http://security.ubuntu.com wily-security/multiverse Sources [2,784 B]  
Get:7 http://security.ubuntu.com wily-security/main amd64 Packages [172 kB]  
Get:8 http://security.ubuntu.com wily-security/restricted amd64 Packages [10.9 kB]  
Get:9 http://security.ubuntu.com wily-security/universe amd64 Packages [56.2 kB]  
Get:10 http://security.ubuntu.com wily-security/multiverse amd64 Packages [6,248 B]  
Get:11 http://security.ubuntu.com wily-security/main i386 Packages [169 kB]  
Get:12 http://security.ubuntu.com wily-security/restricted i386 Packages [10.8 kB]  
100% [Waiting for headers] [Waiting for headers] 73.8 kB/s 0s
```

Step 2 – Install Java and Set JAVA_HOME

//This first thing to do is to setup the webupd8 ppa on your system. Run the following command and proceed.

```
$ sudo apt-add-repository ppa:webupd8team/java
```

```
$ sudo apt-get update
```

//After setting up the ppa repository, update the package cache as well.

//Install the Java 8 installer

```
$ sudo apt-get install oracle-java8-installer
```

// After the installation is finished, Oracle Java is setup. Run the java command again to check the version and vendor.

```
ubuntu@ubuntu-VirtualBox:~$ sudo apt-get install oracle-java8-installer
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  libntdb1 python-ntdb
Use 'apt-get autoremove' to remove them.
The following extra packages will be installed:
  gsfonts-x11 java-common
Suggested packages:
  default-jre equivs binfmt-support visualvm ttf-baekmuk ttf-unfonts
  ttf-unfonts-core ttf-kochi-gothic ttf-sazanami-gothic ttf-kochi-mincho
  ttf-sazanami-mincho ttf-arphic-uming
The following NEW packages will be installed:
  gsfonts-x11 java-common oracle-java8-installer
3 upgraded, 3 newly installed, 0 to remove and 0 not upgraded.
Need to get 163 kB of archives.
After this operation, 511 kB of additional disk space will be used.
Do you want to continue? [Y/n] y
Get:1 http://ppa.launchpad.net/webupd8team/java/ubuntu/ wily/main oracle-java8-i
nstaller all 8u101+8u101arm-1~webupd8~2 [23.6 kB]
```

[or]

```
$ sudo apt-get install default-jdk
```

```
$ java -version
```

```
ubuntu@ubuntu-VirtualBox:~$ java -version
java version "1.8.0_101"
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)
```

Step 3 – Add a dedicated Hadoop user

```
$ sudo addgroup hadoop
```

```
ubuntu@ubuntu-VirtualBox:~$ sudo addgroup hadoop
Adding group `hadoop' (GID 1001) ...
Done.
```

\$ sudo adduser --ingroup hadoop hduser

```
ubuntu@ubuntu-VirtualBox:~$ sudo adduser --ingroup hadoop hduser
Adding user `hduser' ...
Adding new user `hduser' (1001) with group `hadoop' ...
Creating home directory `/home/hduser' ...
Copying files from `/etc/skel' ...
Enter new UNIX password:
Retype new UNIX password:
passwd: password updated successfully
Changing the user information for hduser
Enter the new value, or press ENTER for the default
    Full Name []:
    Room Number []:
    Work Phone []:
    Home Phone []:
    Other []:
Is the information correct? [Y/n] y
```

// Add hduser to sudo user group

\$ sudo adduser hduser sudo

```
ubuntu@ubuntu-VirtualBox:~$ sudo adduser hduser sudo
Adding user `hduser' to group `sudo' ...
Adding user hduser to group sudo
Done.
ubuntu@ubuntu-VirtualBox:~$
```

Step 4 – Install SSH and Create Certificates

\$ sudo apt-get install ssh

```
ubuntu@ubuntu-VirtualBox:~$ sudo apt-get install ssh
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  libntdb1 python-ntdb
Use 'apt-get autoremove' to remove them.
The following extra packages will be installed:
  libck-connector0 ncurses-term openssh-server openssh-sftp-server
  ssh-import-id
Suggested packages:
  rssh molly-guard monkeysphere
The following NEW packages will be installed:
  libck-connector0 ncurses-term openssh-server openssh-sftp-server ssh
  ssh-import-id
0 upgraded, 6 newly installed, 0 to remove and 8 not upgraded.
Need to get 661 kB of archives.
```

\$ su hduser

```
ubuntu@ubuntu-VirtualBox:~$ su hduser
Password:
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.
```

\$ ssh-keygen -t rsa -P ""

```
hduser@ubuntu-VirtualBox:~$ ssh-keygen -t rsa -P ""
Generating public/private rsa key pair.
Enter file in which to save the key (/home/hduser/.ssh/id_rsa):
Created directory '/home/hduser/.ssh'.
Your identification has been saved in /home/hduser/.ssh/id_rsa.
Your public key has been saved in /home/hduser/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:K/F5oNmAqhY02Axp0vzew4EnrN+UGgDTgxIiFPHpT7Q hduser@ubuntu-VirtualBox
The key's randomart image is:
+---[RSA 2048]-----+
|=@o                |
|@.* .              |
|=* * o             |
|.o= *.+            |
|. .=.EooS          |
|...= *B +          |
| o. *+.= .         |
| o o .. .          |
|o                  |
+-----[SHA256]-----+
```

// Set Environmental variables

\$ cat \$HOME/.ssh/id_rsa.pub >> \$HOME/.ssh/authorized_keys

```
hduser@ubuntu-VirtualBox:~$ cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_keys
```

Step 5 – Check if SSH works

\$ ssh localhost


```

hduser@ubuntu-VirtualBox: ~
hduser@ubuntu-VirtualBox:~$ ssh localhost
The authenticity of host 'localhost (127.0.0.1)' can't be established.
ECDSA key fingerprint is SHA256:+j+WF1JPs00vL5mgcc7v9A/rU8jVQEHE8WfLmt2aEo8.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'localhost' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 15.10 (GNU/Linux 4.2.0-41-generic x86_64)

 * Documentation:  https://help.ubuntu.com/

9 packages can be updated.
9 updates are security updates.

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

```

Step 6 – Install Hadoop

```
$ wget https://archive.apache.org/dist/hadoop/core/hadoop-2.8.4/hadoop-2.8.4.tar.gz
```

```
// Extract Hadoop-2.8.4
```

```
$ sudo tar xvzf hadoop-2.8.4.tar.gz
```

```
hduser@ubuntu-VirtualBox:~$ tar xvzf hadoop-2.7.2.tar.gz
```

```
// Create a folder 'hadoop' in /usr/local
```

```
$ sudo mkdir -p /usr/local/hadoop
```

```
hduser@ubuntu-VirtualBox:~$ sudo mkdir -p /usr/local/hadoop
[sudo] password for hduser:
```

```
// Move the Hadoop folder to /usr/local/hadoop
```

```
$ sudo mv hadoop-2.8.4 /usr/local/hadoop
```

```
hduser@ubuntu-VirtualBox:~$ sudo mv hadoop-2.7.2 /usr/local/hadoop
```

```
// Assigning read and write access to Hadoop folder
```

```
$ sudo chown -R hduser:hadoop /usr/local/hadoop
```

```
hduser@ubuntu-VirtualBox:/usr/local/hadoop/hadoop-2.7.2$ sudo chown hduser:hadoop
p -R /usr/local/hadoop
hduser@ubuntu-VirtualBox:/usr/local/hadoop/hadoop-2.7.2$
```

Step 7 - Modify Hadoop config files

//Hadoop Environmental variable setting – The following files will be modified

1. `~/bashrc`
2. `/usr/local/hadoop/hadoop-2.8.4/etc/hadoop/hadoop-env.sh`
3. `/usr/local/hadoop/hadoop-2.8.4/etc/hadoop/core-site.xml`
4. `/usr/local/hadoop/hadoop-2.8.4/etc/hadoop/hdfs-site.xml`
5. `/usr/local/hadoop/hadoop-2.8.4/etc/hadoop/yarn-site.xml`
6. `/usr/local/hadoop/hadoop-2.8.4/etc/hadoop/mapred-site.xml.template`

\$ sudo nano ~/.bashrc

// Add the following lines at the end of the file

```
export JAVA_HOME=/usr/lib/jvm/java-8-oracle
export HADOOP_HOME=/usr/local/hadoop/hadoop-2.8.4
export PATH=$PATH:$HADOOP_HOME/bin
export PATH=$PATH:$HADOOP_HOME/sbin
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export YARN_HOME=$HADOOP_HOME
HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export HADOOP_OPTS="-D.java.library.path=$HADOOP_HOME/lib"
export PATH=$PATH:/usr/local/hadoop/hadoop-2.8.4/bin
```

```
hduser@ubuntu-VirtualBox: ~
GNU nano 2.4.2 File: /home/hduser/.bashrc

if ! shopt -oq posix; then
  if [ -f /usr/share/bash-completion/bash_completion ]; then
    . /usr/share/bash-completion/bash_completion
  elif [ -f /etc/bash_completion ]; then
    . /etc/bash_completion
  fi
fi

#HADOOP VARIABLES START
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
export HADOOP_INSTALL=/usr/local/hadoop/hadoop-2.7.2
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib"
#HADOOP VARIABLES END

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Spell ^_ Go To Line
```

// Configure Hadoop Files

\$ cd /usr/local/hadoop/hadoop-2.8.4/etc/hadoop/

\$ sudo nano hadoop-env.sh

```
hduser@ubuntu-VirtualBox: /usr/local/hadoop/hadoop-2.7.2/etc/hadoop$ cd
hduser@ubuntu-VirtualBox: ~$ cd /usr/local/hadoop/hadoop-2.7.2/etc/hadoop
hduser@ubuntu-VirtualBox: /usr/local/hadoop/hadoop-2.7.2/etc/hadoop$ ls
capacity-scheduler.xml      httpfs-env.sh              mapred-env.sh
configuration.xsl           httpfs-log4j.properties   mapred-queues.xml.template
container-executor.cfg      httpfs-signature.secret   mapred-site.xml.template
core-site.xml               httpfs-site.xml            slaves
hadoop-env.cmd              kms-acls.xml               ssl-client.xml.example
hadoop-env.sh               kms-env.sh                 ssl-server.xml.example
hadoop-metrics2.properties kms-log4j.properties      yarn-env.cmd
hadoop-metrics.properties  kms-site.xml               yarn-env.sh
hadoop-policy.xml           log4j.properties          yarn-site.xml
hdfs-site.xml               mapred-env.cmd
hduser@ubuntu-VirtualBox: /usr/local/hadoop/hadoop-2.7.2/etc/hadoop$ sudo nano ha
dooop-env.sh
```

// Add following line in hadoop-env.sh – Set JAVA variable in Hadoop

```
# The java implementation to use.
export JAVA_HOME=/usr/lib/jvm/java-8-oracle
```

```
hduser@ubuntu-VirtualBox: /usr/local/hadoop/hadoop-2.7.2/etc/hadoop
GNU nano 2.4.2 File: hadoop-env.sh Modified
# set JAVA_HOME in this file, so that it is correctly defined on
# remote nodes.
# The java implementation to use.
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
# The jsvc implementation to use. Jsvc is required to run secure datanodes
# that bind to privileged ports to provide authentication of data transfer
# protocol. Jsvc is not required if SASL is configured for authentication of
# data transfer protocol using non-privileged ports.
#export JSVC_HOME=${JSVC_HOME}
export HADOOP_CONF_DIR=${HADOOP_CONF_DIR:-"/etc/hadoop"}
# Extra Java CLASSPATH elements. Automatically insert capacity-scheduler.
for f in $HADOOP_HOME/contrib/capacity-scheduler/*.jar; do
  if [ "$HADOOP_CLASSPATH" ]; then
    export HADOOP_CLASSPATH=$HADOOP_CLASSPATH:$f
  else

```

// Create datanode and namenode

\$ sudo mkdir -p /usr/local/hadoop_tmp/hdfs/namenode

\$ sudo mkdir -p /usr/local/hadoop_tmp/hdfs/datanode

// Changing ownership to hadoop_tmp

\$ sudo chown -R hduser:hadoop /usr/local/hadoop_tmp

```
hduser@ubuntu-VirtualBox:~$ sudo mkdir -p /usr/local/hadoop_tmp/hdfs/namenode
hduser@ubuntu-VirtualBox:~$ sudo mkdir -p /usr/local/hadoop_tmp/hdfs/datanode
hduser@ubuntu-VirtualBox:~$ sudo chown -R hduser:hadoop -R /usr/local/hadoop_tmp
```

// Edit hdfs-site.xml

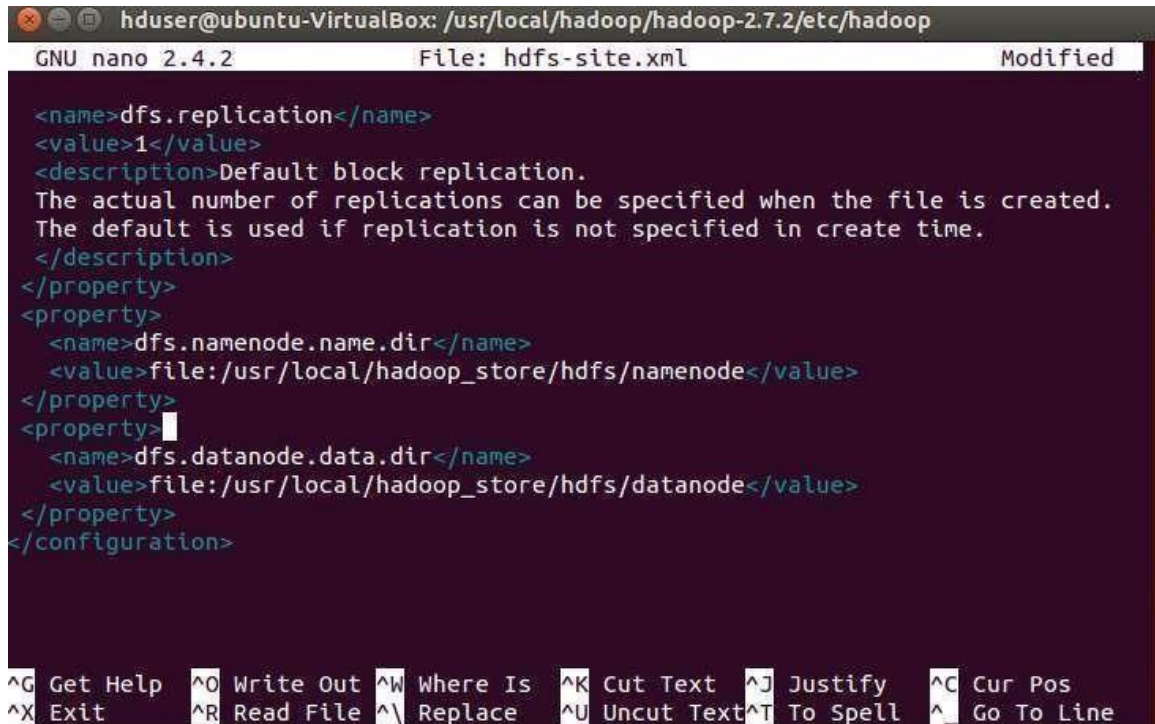
\$ sudo nano hdfs-site.xml

// Add the following lines between <configuration> </configuration>

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>file:/usr/local/hadoop_tmp/hdfs/namenode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
```



```
<value>file:/usr/local/hadoop_tmp/hdfs/datanode</value>
</property>
</configuration>
```



```
hduser@ubuntu-VirtualBox: /usr/local/hadoop/hadoop-2.7.2/etc/hadoop
GNU nano 2.4.2 File: hdfs-site.xml Modified

<name>dfs.replication</name>
<value>1</value>
<description>Default block replication.
The actual number of replications can be specified when the file is created.
The default is used if replication is not specified in create time.
</description>
</property>
<property>
  <name>dfs.namenode.name.dir</name>
  <value>file:/usr/local/hadoop_store/hdfs/namenode</value>
</property>
<property>
  <name>dfs.datanode.data.dir</name>
  <value>file:/usr/local/hadoop_store/hdfs/datanode</value>
</property>
</configuration>

^G Get Help  ^O Write Out  ^W Where Is  ^K Cut Text  ^J Justify   ^C Cur Pos
^X Exit      ^R Read File  ^\ Replace   ^U Uncut Text ^T To Spell  ^_ Go To Line
```

// Edit core-site.xml

\$ sudo nano core-site.xml

// Add the following lines between <configuration> </configuration>

```
<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
```

// Edit yarn-site.xml

\$ sudo nano yarn-site.xml

// Add the following lines between <configuration> </configuration>

```
<configuration>
<property>
```

```

<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
<value>org.apache.hadoop.mapred.Shuffle-Handler</value>
</property>
</configuration>

```

// Edit mapred-site.xml

```

$ cp /usr/local/hadoop/hadoop-2.8.4/etc/hadoop/mapred-site.xml.template
/usr/local/hadoop/hadoop-2.8.4/etc/hadoop/mapred-site.xml

```

```

hduser@ubuntu-VirtualBox:~$ cp /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/mapred-
site.xml.template /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/mapred-site.xml

```

\$ sudo nano mapred-site.xml

// Add the following lines between <configuration> </configuration>

```

<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>

```

Step 8 – Format Hadoop File System

```
$ cd /usr/local/hadoop/hadoop-2.8.4/bin
```

\$ hadoop namenode -format

```

hduser@ubuntu-VirtualBox:/usr/local/hadoop$ hadoop namenode -format
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

16/07/15 22:50:27 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode

```

Step 9 - Start Hadoop

```
$ cd /usr/local/hadoop/hadoop-2.8.4/sbin
```

// Starting dfs services

\$ start-dfs.sh

```

hduser@ubuntu-VirtualBox: /usr/local/hadoop/hadoop-2.7.2/sbin$ start-dfs.sh
16/07/15 22:55:47 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/hadoop/hadoop-2.7.2/logs/hadoop-hduser-namenode-ubuntu-VirtualBox.out
localhost: starting datanode, logging to /usr/local/hadoop/hadoop-2.7.2/logs/hadoop-hduser-datanode-ubuntu-VirtualBox.out
Starting secondary namenodes [0.0.0.0]
The authenticity of host '0.0.0.0 (0.0.0.0)' can't be established.
ECDSA key fingerprint is SHA256:+j+WF1JPs00vL5mgcc7v9A/rU8jVQEHE8WfLmt2aEo8.
Are you sure you want to continue connecting (yes/no)? yes
0.0.0.0: Warning: Permanently added '0.0.0.0' (ECDSA) to the list of known hosts
.
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/hadoop-2.7.2/logs/hadoop-hduser-secondarynamenode-ubuntu-VirtualBox.out

```

// Starting mapreduce services

\$ start-yarn.sh

```

hduser@ubuntu-VirtualBox: /usr/local/hadoop/hadoop-2.7.2/sbin$ start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/hadoop-2.7.2/logs/yarn-hduser-resourcemanager-ubuntu-VirtualBox.out
localhost: starting nodemanager, logging to /usr/local/hadoop/hadoop-2.7.2/logs/yarn-hduser-nodemanager-ubuntu-VirtualBox.out

```

\$ jps

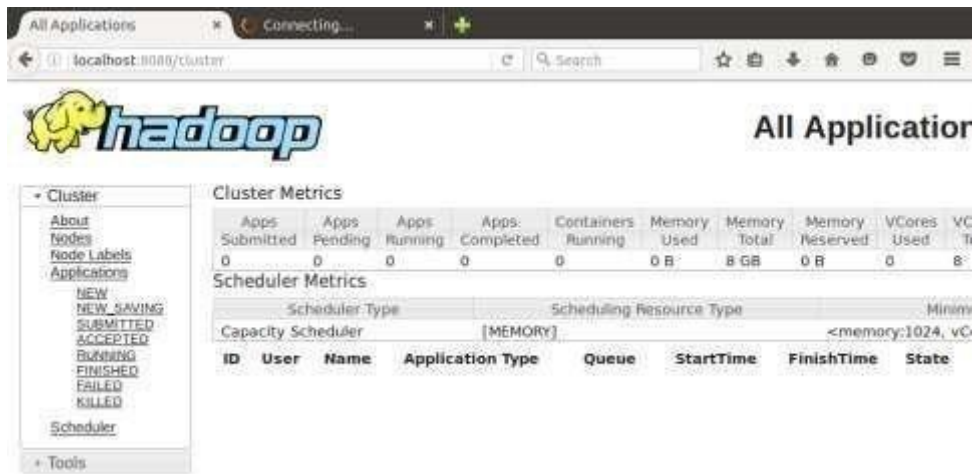
```

hduser@ubuntu-VirtualBox: /usr/local/hadoop/hadoop-2.7.2/sbin$ jps
12425 SecondaryNameNode
12609 ResourceManager
12733 NodeManager
13131 Jps
12205 DataNode
12080 NameNode

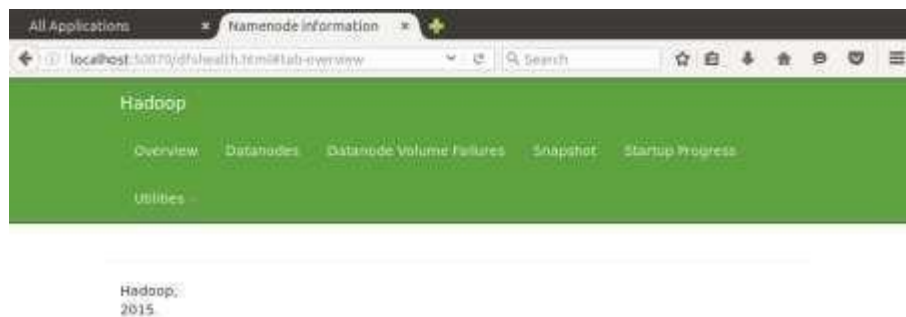
```

Step 10 - Check Hadoop through web UI

Go to browser type <http://localhost:8088> – All Applications Hadoop Cluster



Go to browser type <http://localhost:50070> – Hadoop Namenode



Step 11 - Stop Hadoop

\$ stop-dfs.sh

\$ stop-yarn.sh

RESULT

Thus the procedure to install single-node Hadoop is executed successfully.

EX. NO:

DATE:

DEMONSTRATION OF MAP AND REDUCE TASKS

AIM

To write a wordcount program to demonstrate the use of Map and Reduce tasks.


THEORY

Hadoop MapReduce is a programming paradigm at the heart of Apache Hadoop for providing massive scalability across hundreds or thousands of Hadoop clusters on commodity hardware. The MapReduce model processes large unstructured data sets with a distributed algorithm on a Hadoop cluster.

The term MapReduce represents two separate and distinct tasks Hadoop programs perform-Map Job and Reduce Job. Map job scales takes data sets as input and processes them to produce key value pairs. Reduce job takes the output of the Map job i.e. the key value pairs and aggregates them to produce desired results. The input and output of the map and reduce jobs are stored in HDFS.

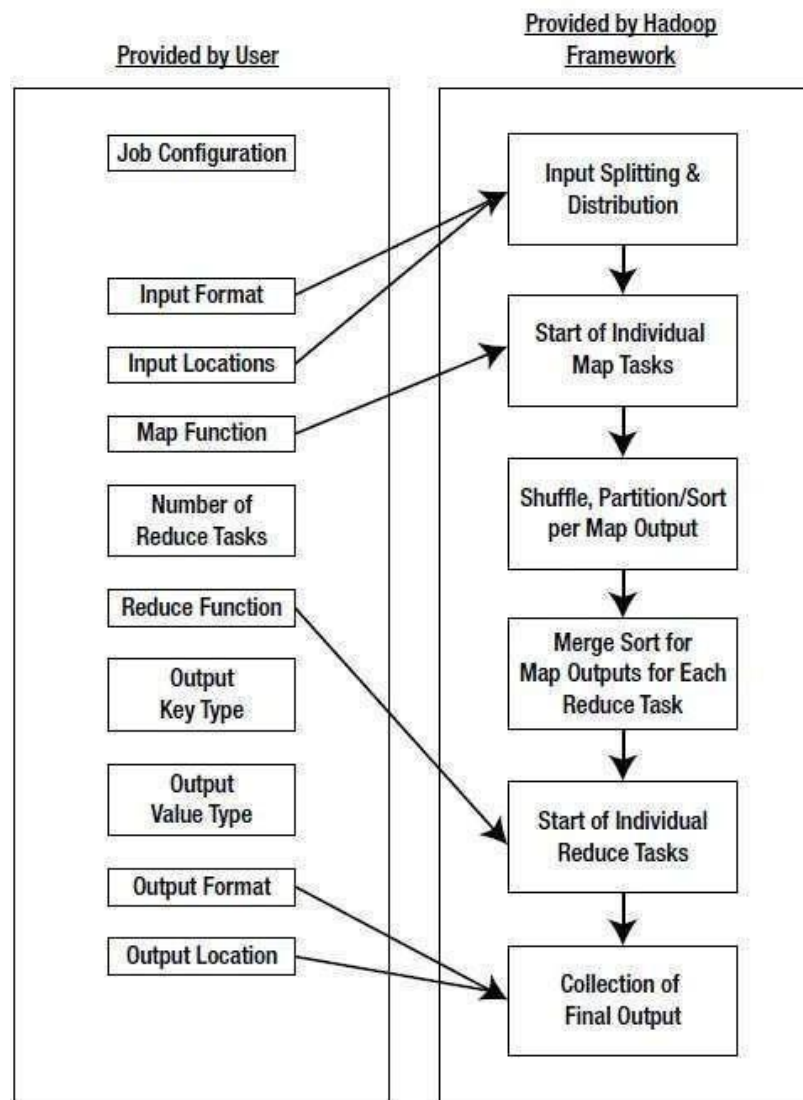
MAPREDUCE

Hadoop MapReduce (Hadoop Map/Reduce) is a software framework for distributed processing of large data sets on computing clusters. It is a sub-project of the Apache Hadoop project. Apache Hadoop is an open-source framework that allows to store and process big data in a distributed environment across clusters of computers using simple programming models. MapReduce is the core component for data processing in Hadoop framework. In layman's term Mapreduce helps to split the input data set into a number of parts and run a program on all data parts parallel at once. The term MapReduce refers to two separate and distinct tasks. The first is the map operation, takes a set of data and converts it into another set of data, where individual elements are broken down into tuples (key/value pairs). The reduce operation combines those data tuples based on the key and accordingly modifies the value of the key.



MAPREDUCE ARCHITECTURE

The figure shown below illustrates the various parameters and modules that can be configured during a MapReduce operation:



MapReduce Architecture

JobConf is the framework used to provide various parameters of a MapReduce job to the Hadoop for execution. The Hadoop platforms executes the programs based on configuration set using JobConf. The parameters being Map Function, Reduce Function, combiner , Partitioning function, Input and Output format. Partitioner controls the shuffling of the tuples when being sent from Mapper node to Reducer nodes. The total number of partitions done in the tuples is equal to the number of reduce nodes. In simple terms based on the function output the tuples are transmitted through different reduce nodes.

Input Format describes the format of the input data for a MapReduce job. Input location specifies the location of the datafile. Map Function/ Mapper convert the data into key value pair. For example let's consider daily temperature data of 100 cities for the past 10 years. In this the map function is written such a way that every temperature being mapped to the corresponding

city. Reduce Function reduces the set of tuples which share a key to a single tuple with a change in the value. In this example if we have to find the highest temperature recorded in the city the reducer function is written in such a way that it return the tuple with highest value i.e: highest temperature in the city in that sample data.

The number of Map and Reduce nodes can also be defined. You can set Partitioner function which partitions and transfer the tuples which by default is based on a hash function. In other words we can set the options such that a specific set of key value pairs are transferred to a specific reduce task. For example if your key value consists of the year it was recorded, then we can set the parameters such that all the keys of specific year are transferred to a same reduce task. The Hadoop framework consists of a single master and many slaves. Each master has JobTracker and each slave has TaskTracker. Master distributes the program and data to the slaves. Task tracker, as the name suggests keep track of the task directed to it and relays the information to the JobTracker. The JobTracker monitors all the status reports and re-initiates the failed tasks if any.

Combiner classes are run on map task nodes. It takes the tuples emitted by Map nodes as input. It basically does reduce operation on the tuples emitted by the map node. It is like a pre-reduce task to save a lot of bandwidth. We can also pass global constants to all the nodes using „Counters’. They can be used to keep track of all events in map and reduce tasks. For example we can pass a counter to calculate the statistics of an event beyond a certain threshold.

PROCEDURE

Step 1 - Open Terminal

```
$ su hduser
```

Password:

Step 2 - Start dfs and mapreduce services

```
$ cd /usr/local/hadoop/hadoop-2.8.4/sbin
```

```
$ start-dfs.sh
```

```
$ start-yarn.sh
```

```
$ jps
```

Step 3 - Check Hadoop through web UI

```
// Go to browser type http://localhost:8088 – All Applications Hadoop Cluster
```

```
// Go to browser type http://localhost:50070 – Hadoop Namenode
```

Step 4 – Open New Terminal

```
$ cd Desktop/
```

```
$ mkdir inputdata
```

```
$ cd inputdata/
```

```
$ echo "Hai, Hello, How are you? How is your health?" >> hello.txt
```

```
$ cat >> hello.txt
```

Step 5 – Go back to old Terminal

```
$ hadoop fs -mkdir /folder
```

```
$ hadoop fs -copyFromLocal /home/hduser/Desktop/inputdata/hello.txt /folder
```

// Check in hello.txt in Namenode using Web UI

Step 6 – Download and open eclipse by creating workspace

Create a new java project.

Step 7 – Add jar to the project

You need to remove dependencies by adding jar files in the hadoop source folder. Now Click on **Project** tab and go to Properties. Under Libraries tab, click Add External JARs and select all the jars in the folder (click on 1st jar, and Press Shift and Click on last jar to select all jars in between and click ok)

/usr/local/hadoop/hadoop-2.8.4/share/hadoop/common and

/usr/local/hadoop/hadoop-2.8.4/share/hadoop/mapreduce folders.

Step -8 – WordCount Program

Create 3 java files named

- **WordCount.java**
- **WordCountMapper.java**
- **WordCountReducer.java**

WordCount.java

```
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;

import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
import org.apache.hadoop.io.Text;

public class WordCount extends Configured implements Tool {

    @Override
    public int run(String[] arg0) throws Exception {
        // TODO Auto-generated method stub
        if(arg0.length<2)
```



```

        {
            System.out.println("check the command line arguments");
        }
        JobConf conf=new JobConf(WordCount.class);
        FileInputFormat.setInputPaths(conf, new Path(arg0[0]));
        FileOutputFormat.setOutputPath(conf, new Path(arg0[1]));
        conf.setMapperClass(WordMapper.class);
        conf.setReducerClass(WordReducer.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        JobClient.runJob(conf);

        return 0;
    }
    public static void main(String args[]) throws Exception
    {
        int exitcode=ToolRunner.run(new WordCount(), args);
        System.exit(exitcode);
    }
}

```

WordCountMapper.java

```

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.Mapper;

public class WordCountMapper extends MapReduceBase implements
Mapper<LongWritable,Text,Text,IntWritable>
{
    @Override
    public void map(LongWritable arg0, Text arg1, OutputCollector<Text, IntWritable>
arg2, Reporter arg3)
        throws IOException {
        // TODO Auto-generated method stub

        String s=arg1.toString();
        for(String word:s.split(" "))
        {
            arg2.collect(new Text(word),new IntWritable(1));
        }
    }
}

```

WordCountReducer.java

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.io.Text;

public class WordCountReducer implements Reducer<Text,IntWritable,Text,IntWritable> {
    @Override
    public void configure(JobConf arg0) {
        // TODO Auto-generated method stub
    }
    @Override
    public void close() throws IOException {
        // TODO Auto-generated method stub
    }
    @Override
    public void reduce(Text arg0, Iterator<IntWritable> arg1, OutputCollector<Text,
IntWritable> arg2, Reporter arg3)
        throws IOException {
        // TODO Auto-generated method stub
        int count=0;
        while(arg1.hasNext())
        {
            IntWritable i=arg1.next();
            count+=i.get();
        }
        arg2.collect(arg0,new IntWritable(count));
    }
}
```

Step 9 - Creatr JAR file

Now Click on the Run tab and click Run-Configurations. Click on New Configuration button on the left-top side and Apply after filling the following properties.

Step 10 - Export JAR file

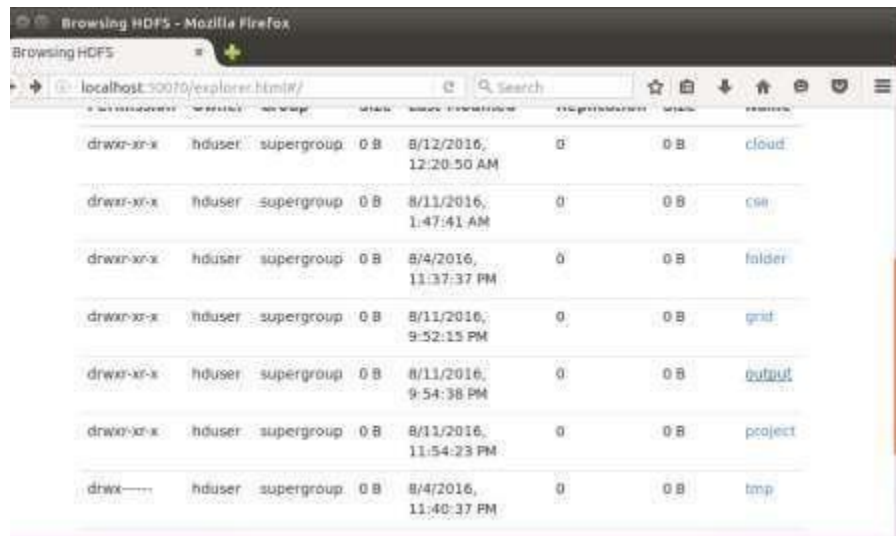
Now click on File tab and select Export. under Java, select Runnable Jar.

In Launch Config – select the config fie you created in **Step 9** (WordCountConfig).

- ☐ Select an export destination (lets say desktop.)
- ☐ Under Library handling, select Extract Required Libraries into generated JAR and click Finish.
- ☐ Right-Click the jar file, go to Properties and under **Permissionstab**, Check Allow executing file as a program. and give Read and Write access to all theusers

Step 11 – Go back to old Terminal for Execution of WordCount Program

```
$hadoop jar wordcount.jar/usr/local/hadoop/input/usr/local/hadoop/output
```

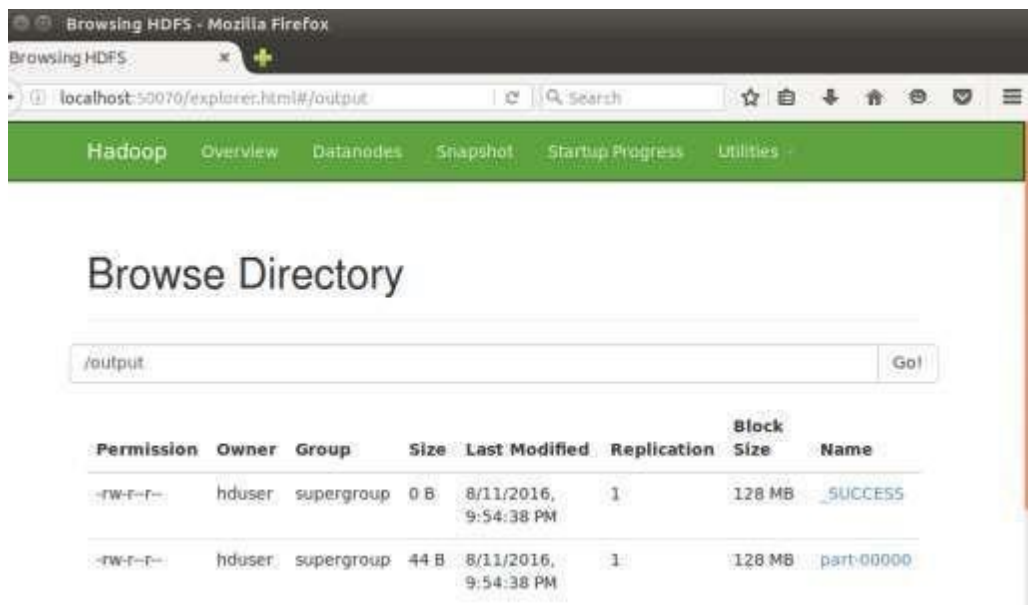


A screenshot of a web browser showing the HDFS Explorer interface. The browser address bar shows 'localhost:50070/explorer.html#/'. The interface displays a list of files and folders in HDFS. The table has columns for permissions, owner, group, size, last modified, replication, and name. The files listed are 'cloud', 'cse', 'folder', 'grid', 'output', 'project', and 'tmp'.

Permissions	Owner	Group	Size	Last Modified	Replication	Name
drwxr-xr-x	hduser	supergroup	0 B	8/12/2016, 12:20:50 AM	0	cloud
drwxr-xr-x	hduser	supergroup	0 B	8/11/2016, 1:47:41 AM	0	cse
drwxr-xr-x	hduser	supergroup	0 B	8/4/2016, 11:37:37 PM	0	folder
drwxr-xr-x	hduser	supergroup	0 B	8/11/2016, 9:52:15 PM	0	grid
drwxr-xr-x	hduser	supergroup	0 B	8/11/2016, 9:54:38 PM	0	output
drwxr-xr-x	hduser	supergroup	0 B	8/11/2016, 11:54:23 PM	0	project
drwxr-xr-x	hduser	supergroup	0 B	8/4/2016, 11:40:37 PM	0	tmp

Step 12 – To view results in old Terminal

```
$hdfs dfs -cat /usr/local/hadoop/output/part-r-00000
```



A screenshot of the HDFS Explorer interface showing the contents of the '/output' directory. The browser address bar shows 'localhost:50070/explorer.html#/output'. The interface displays a table with columns for Permission, Owner, Group, Size, Last Modified, Replication, Block Size, and Name. The files listed are '_SUCCESS' and 'part-00000'.

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	hduser	supergroup	0 B	8/11/2016, 9:54:38 PM	1	128 MB	_SUCCESS
-rw-r--r--	hduser	supergroup	44 B	8/11/2016, 9:54:38 PM	1	128 MB	part-00000

Step 13 - To Remove folders created using hdfs

```
$ hdfs dfs -rm -R /usr/local/hadoop/output
```

RESULT

Thus the program to write a wordcount program to that demonstrates the use of Map and Reduce tasks is done and output is verified successfully.

EX. NO:

DATE:

IMPLEMENTATION OF MAX TEMPERATURE MAPREDUCE PROGRAM

AIM

To implement the Max temperature MapReduce program to identify the year wise maximum temperature from the sensor data.

Description

Sensors sense weather data in big text format containing station ID, year, date, time, temperature, quality etc. from each sensor and store it in single line. Suppose thousands of data sensors are there, then we have thousands of records with no particular order. We require only year and maximum temperature of particular quality in that year.

For example:

Input string from sensor:

002902907099999**1902**010720004+64333+023450FM-12+
000599999V0202501N0278199999999N0000001N9-**00331**+
99999098351ADDGF10299199999999999999999999

Here: **1902** is year

0033 is temperature

1 is measurement quality (Range between 0 or 1 or 4 or 5 or 9)

Here each mapper takes input **key** as "byte offset of line" and **value** as "one weather sensor read i.e one line". and parse each line and produce **intermediate key** is "year" and **intermediate value** as "**temperature of certain measurement qualities**" for that year.

The combiner will form set values of temperature. Year and set of values of temperatures is given as input **<key, value>** to reducer and Reducer will produce year and maximum temperature for that year from set of temperature values.

PROGRAM

```
*/  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Job;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
  
import java.io.IOException;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;  
import org.apache.hadoop.mapreduce.Reducer;  
  
//Mapper class
```

```

class MaxTemperatureMapper
extends Mapper<LongWritable, Text, Text, IntWritable> {

    private static final int MISSING = 9999;

    @Override
    public void map(LongWritable key, Text value, Context context)
    throws IOException, InterruptedException {

        String line = value.toString();
        String year = line.substring(15, 19);
        int airTemperature;
        if (line.charAt(87) == '+') { // parseInt doesn't like leading plus signs
            airTemperature = Integer.parseInt(line.substring(88, 92));
        } else {
            airTemperature = Integer.parseInt(line.substring(87, 92));
        }
        String quality = line.substring(92, 93);
        if (airTemperature != MISSING && quality.matches("[01459]")) {
            context.write(new Text(year), new IntWritable(airTemperature));
        }
    }
}

//Reducer class
class MaxTemperatureReducer
extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
    Context context)
    throws IOException, InterruptedException {

        int maxVal = Integer.MIN_VALUE;
        for (IntWritable value : values) {
            maxVal = Math.max(maxVal, value.get());
        }
        context.write(key, new IntWritable(maxVal));
    }
}

//Driver Class

public class MaxTemperature {

    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Usage: MaxTemperature <input path=""> <output path="">");
            System.exit(-1);
        }

        Job job = Job.getInstance(new Configuration());
        job.setJarByClass(MaxTemperature.class);
        job.setJobName("Max temperature");

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
    }
}

```

```
job.setMapperClass(MaxTemperatureMapper.class);  
job.setReducerClass(MaxTemperatureReducer.class);  
  
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);  
  
job.submit();  
}  
}
```

OUTPUT

Input for String :

002902907099999**1902**010720004+64333+023450FM-12+
000599999V0202501N0278199999999N0000001N9-**00331**+
99999098351ADDGF10299199999999999999999999

Output

Output Text contain year and maximum temperature in that year as **1902 33**

RESULT

Thus the implementation of the Max temperature MapReduce program to identify the year wise maximum temperature from the sensor data has been done successfully.